



Applied Deep Learning

Dr. Philippe Blaettchen
Bayes Business School (formerly Cass)

www.bayes.city.ac.uk

Learning objectives of today

Goals: Be ready for your group assignment

How will we do this?

- We will learn and practice another way to tune hyperparameters
- Then, we will take a look at autoencoders, a specific type of neural network
- Finally, we will discuss the case underlying the assignment, as well as some of the key challenges



BAYES
BUSINESS SCHOOL
CITY, UNIVERSITY OF LONDON



Hyperparameter tuning with Keras tuner

Let's try it together in Python



BAYES
BUSINESS SCHOOL
CITY, UNIVERSITY OF LONDON

Hyperparameter tuning process with Keras Tuner

1. Define a function that, given a hyperparameter-setter, creates a model
Within that function, using the hyperparameter-setter, we define the hyperparameter space
2. Define an instance of the Keras Tuner, specifying the type of hyperparameter search
Can use RandomSearch, Hyperband, Sklearn, BayesianOptimization

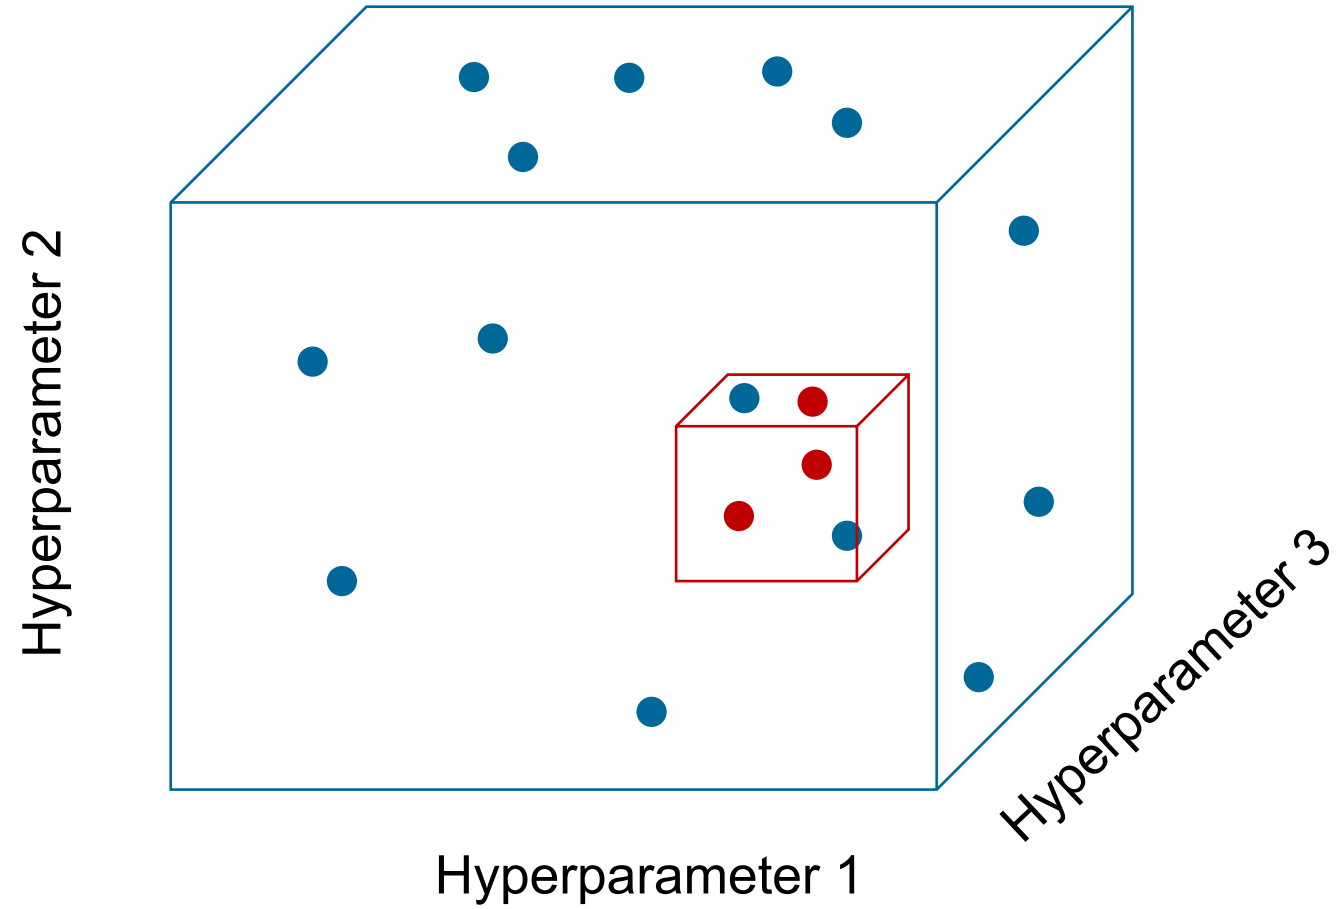


Hyperparameter tuning process with Keras Tuner

1. Define a function that, given a hyperparameter-setter, creates a model
Within that function, using the hyperparameter-setter, we define the hyperparameter space
2. Define an instance of the Keras Tuner, specifying the type of hyperparameter search
Can use RandomSearch, **Hyperband**, Sklearn, BayesianOptimization
3. Let the tuner do its magic
The hyperparameter-setter will automatically choose the “correct” hyperparameters
4. Based on the best parameters found, generate a model, train it, and evaluate it
5. Look at the outcome, possibly search in a smaller grid



Hyperparameter space



Autoencoders

What is an autoencoder?

- A neural network that predicts its own inputs
 - So that we can learn a (compact) representation of the data



Recall that a neural network learns representations



x

Learn $f(\cdot)$ →

$f(x)$

Learn $g(\cdot)$ →

$g(f(x)) \approx y$

E.g., $y = 1$, if it's a cat
 $y = 0$, if it's a BA student

Recall that a neural network learns representations



x

Learn $f(\cdot)$ →

$f(x)$

Learn $g(\cdot)$ →

$g(f(x)) \approx x$




BAYES
BUSINESS SCHOOL
CITY, UNIVERSITY OF LONDON


Why do we want to “copy” the input?

- We don't care about the copy itself (which should be good, nevertheless)
- What we care about is a (compact) representation of the data, $f(x)$ that is good enough to recreate x

The structure of an autoencoder




NOT the structure of an autoencoder



BAYES
BUSINESS SCHOOL
CITY, UNIVERSITY OF LONDON

Let's try it together in Python



BAYES
BUSINESS SCHOOL
CITY, UNIVERSITY OF LONDON

What to use it for?

- Dimensionality reduction (“advanced PCA”)
- **Denoising: train to “recover” non-noisy data from noisy data**
- Anomaly detection: train to represent normal data. When data cannot be predicted well, it is likely to be “anormal”
- Generate new content (such as images): variational autoencoders



BAYES
BUSINESS SCHOOL
CITY, UNIVERSITY OF LONDON

The process for denoising

- We create artificial noise on our data
- We build an autoencoder that takes the noisy data as input, and tries to build an accurate representation of the original
 - To do so with images, we need convolutional layers. Don't worry about how they work, we will get to them soon. You find all the code on using them in the notebook!
 - Training the autoencoder may take quite a bit of time! (I suggest first getting to the training part, then taking a break in the meantime)
- We then can run the autoencoder on new (noisy) data, to create non-noisy versions

Try it out in Python



BAYES
BUSINESS SCHOOL
CITY, UNIVERSITY OF LONDON



An overview of the group assignment

The group assignment


- “To Catch a Thief”: A case about “Shift”, an InsureTech startup that helps insurance companies uncover fraudulent claims
 - In particular, their algorithm flags suspicious cases for claim handlers to investigate
- We have a dataset of vehicle insurance claims. One issue is that it is “unbalanced”
 - only about 1% of cases are frauds
 - What would be the accuracy of an algorithm that always predicts “no fraud”?
- Your tasks:
 - Pre-processing
 - Understanding the relevant metrics
 - Trialing and comparing models of varying complexity
 - Exploring autoencoders as a tool to detect anomalous data
 - Discussing transparency implications of neural networks





Anomaly detection with autoencoders

Outliers and anomalies



Outliers:


- Data points that are distinctly different from other data points
- Can be caused by unavoidable random errors or by systematic errors relating to how data was sampled

Anomalies:


- Outliers or other values that are not expected to exist
- Can be context- or pattern-based:
 - Context: exceptionally high credit card spending on Black Friday versus near-simultaneous spending in New York and London
 - Pattern: high credit card spending every Saturday versus high spending on a day where spending is low in other weeks



Differentiating anomalies from normal observations



Differentiating anomalies from normal observations



What are possible anomalies and how would we detect them?

Consider the following situations:

- A machine produces thousands of screws per minute, every few days the type of screw is changed
- A software developer for a bank downloads a large number of entries from a customer database
- An intermediary supplies fair trade coffee beans

What is the expected outcome in each case?

What is an anomalous outcome?

What data do we observe?



BAYES
BUSINESS SCHOOL
CITY, UNIVERSITY OF LONDON

Detecting anomalies

Supervised anomaly detection:

- A fancy way of saying classification – learn to differentiate between two classes
- We can use the standard toolbox
- When feasible, usually the most failsafe method
- Only works if we know how anormal data looks like, i.e., we have enough data points

Semi-supervised anomaly detection:

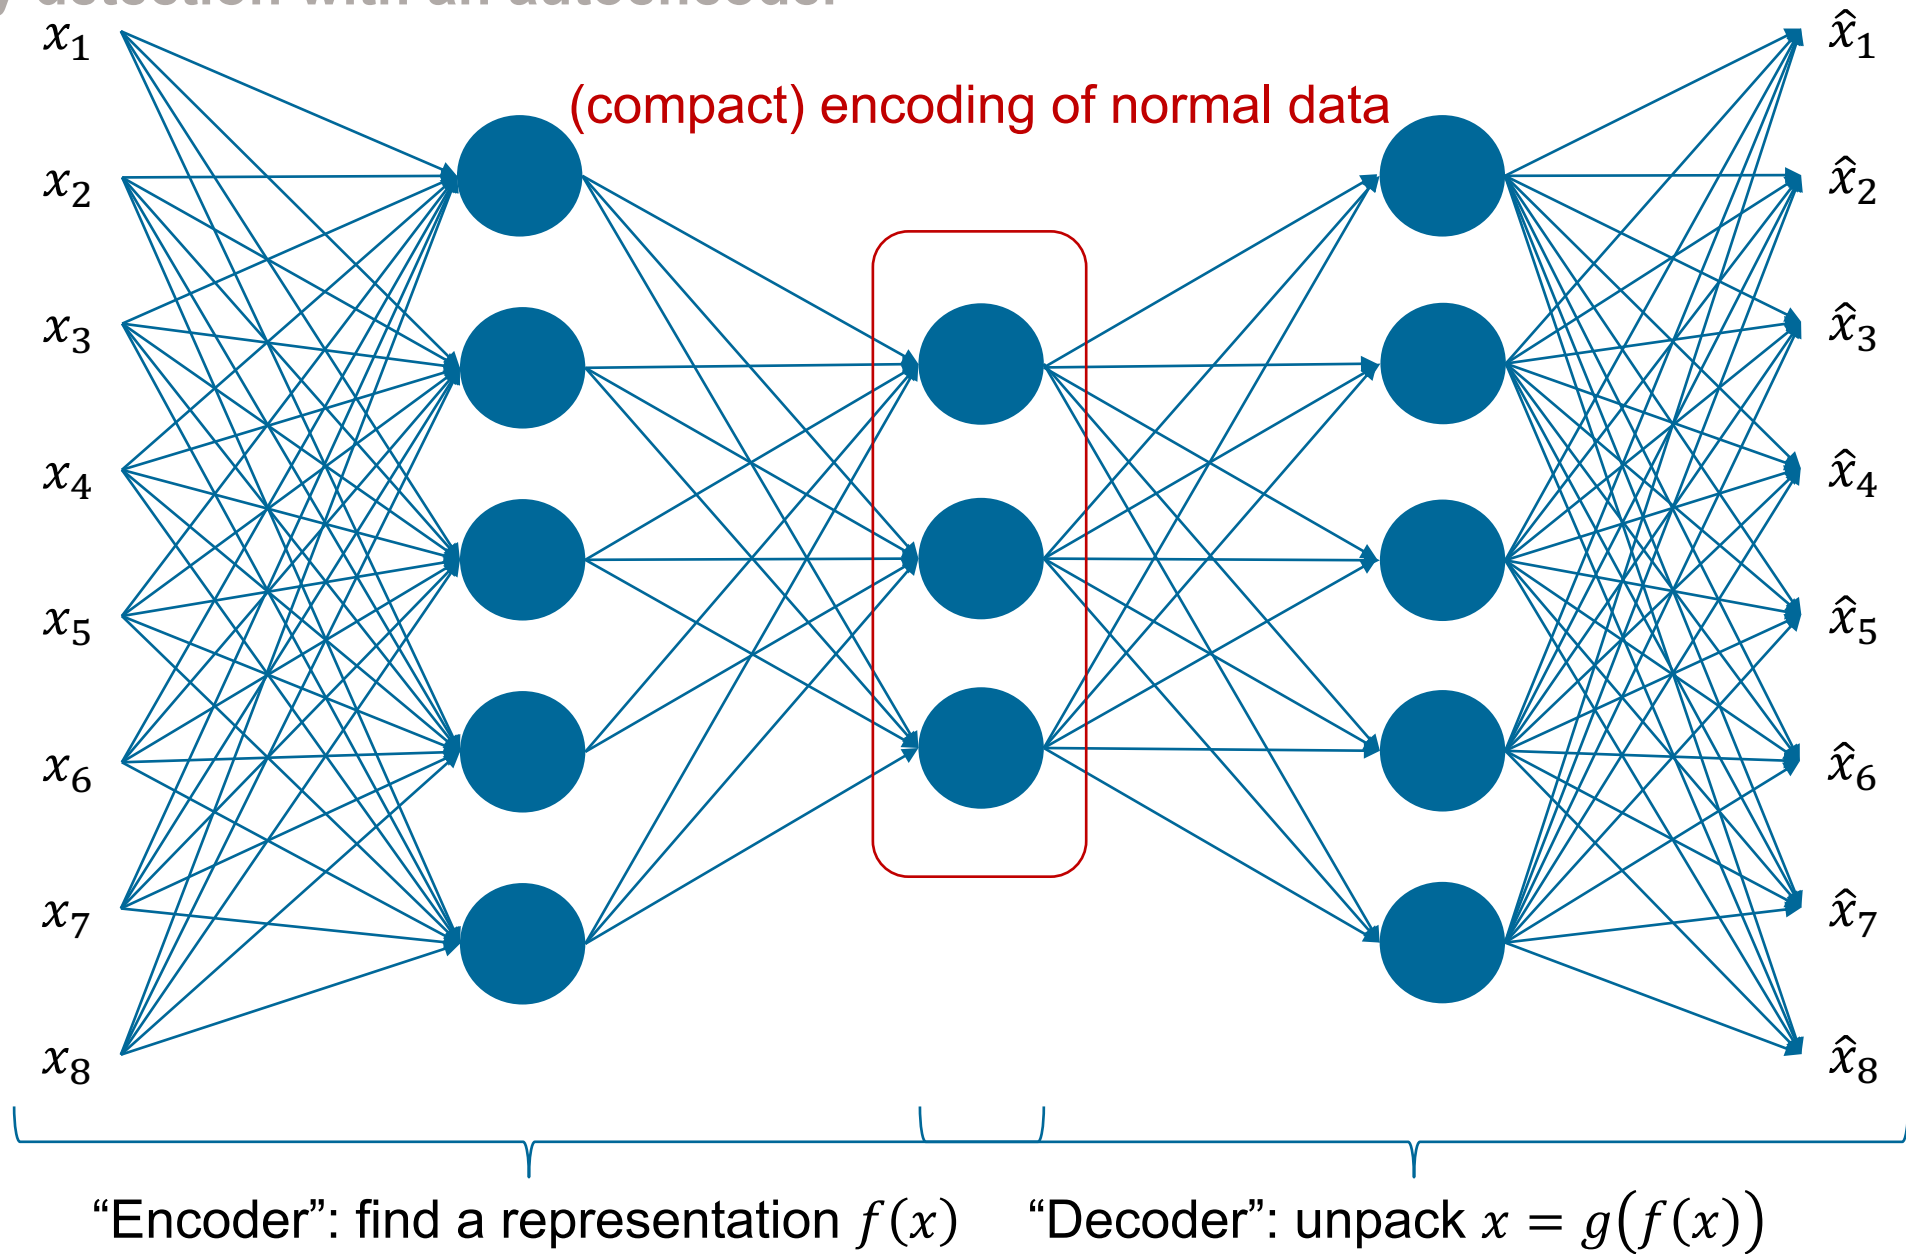
- Learn an efficient representation of normal data and then try to apply this to new data coming in
- We can use autoencoders and other tools
- We don't need to know how anormal data looks like
- Still need to be sure that our normal data is actually normal

Unsupervised anomaly detection:

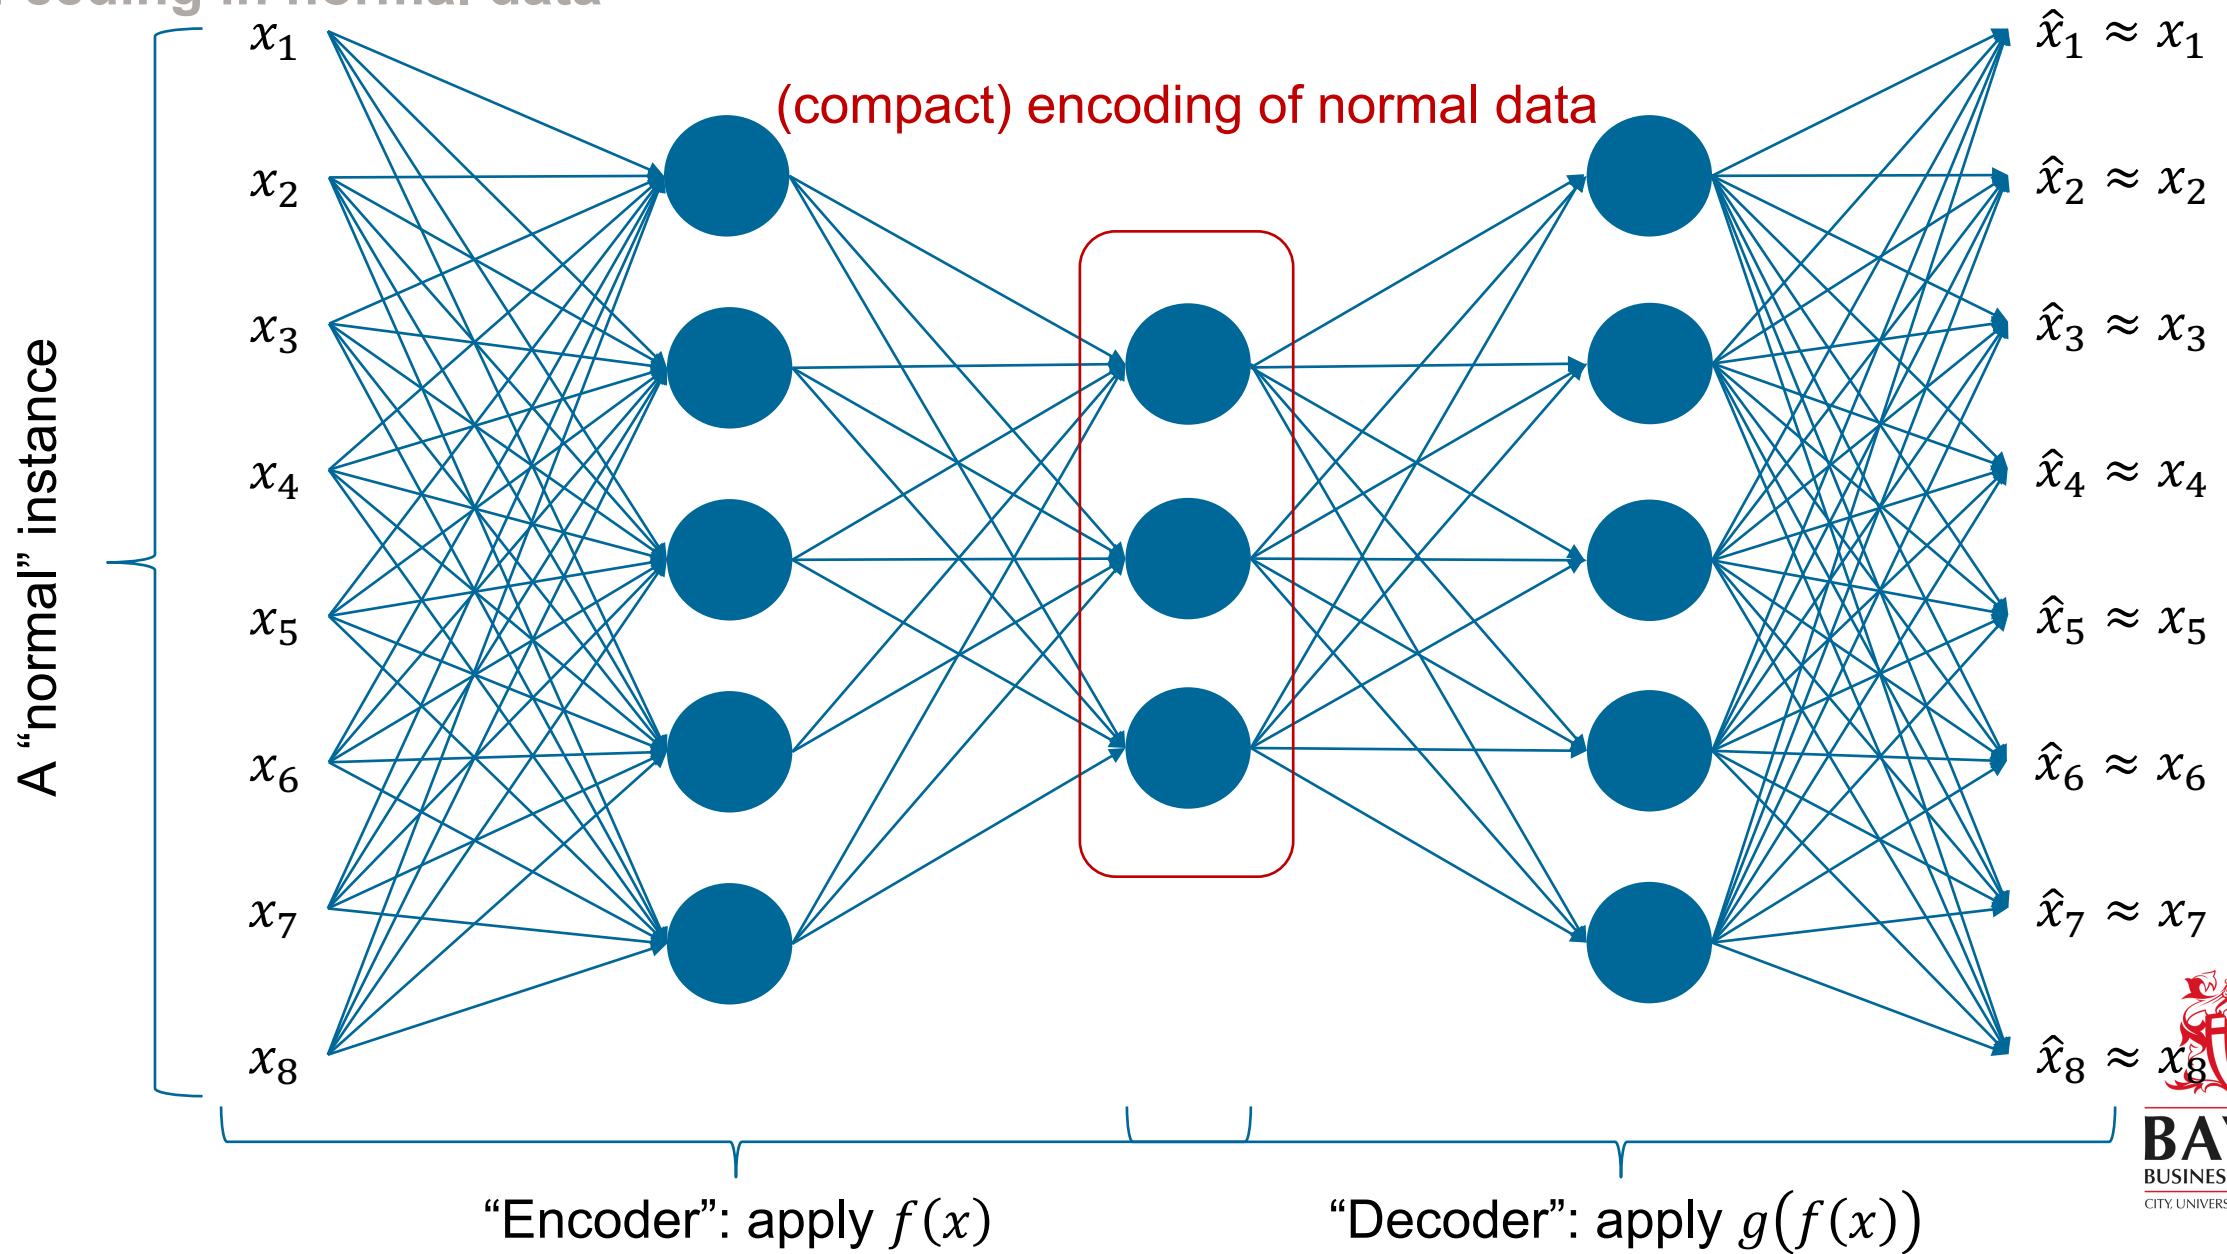
- Learn “how far” datapoints are from each other and recognize the ones that are far away from anything else
- We can use isolation forests and other tools
- We can work with any kind of data
- We don't have many guarantees



Anomaly detection with an autoencoder




Feeding in normal data




Feeding in anormal data

An “anomalous” instance



What should we be observing?





Good luck with the assignment!