



# Applied Deep Learning

Dr. Philippe Blaettchen  
Bayes Business School (formerly Cass)

[www.bayes.city.ac.uk](http://www.bayes.city.ac.uk)

## Learning objectives of today

### **Goals:** Creating neural networks with Keras and TensorFlow

- Understanding the role that frameworks such as TensorFlow play in enabling the design, training, and use of arbitrarily complex neural networks
- Getting started on applying TensorFlow to create different types of feed-forward networks


### **How will we do this?**

- We start with a quick recap of the functioning of neural networks
- We then introduce TensorFlow and Keras as programming frameworks for neural networks
- We create a number of neural networks, using TensorFlow's Sequential API




Recap

# Components




... ..



Source: Liang

## Hidden layers

$$“x” = \mathbf{a}^{[l-1]} \quad z = \mathbf{a}^{[l-1]}\mathbf{w} + b \quad f(z)$$




- $f$  is what we call an “activation function”
- There are many activation functions, and new ones are invented all the time
- Many of these functions do just fine, or slightly better than existing ones

Source: Liang


## Typical activation functions: Rectified Linear Unit

Rectified Linear Unit (ReLU)



$$f(z) = \max\{0, z\}$$

“Derivative”




$$f'(z) = \begin{cases} 0, & \text{if } z < 0 \\ 1, & \text{if } z > 0 \end{cases}$$



## Binary classification

- Input:  $\mathbf{a}^{[L-1]}(i)$
- As usual, we make a linear transformation:  
$$z^{[L]}(i) = \mathbf{a}^{[L-1]}(i) \mathbf{w}^{[L]} + b^{[L]}$$
- We then use the logistic sigmoid function  
$$\hat{y}^{(i)} = f(z^{[L]}(i)) = \sigma(z^{[L]}(i)) = \frac{1}{1 + e^{-z^{[L]}(i)}}$$
- We can interpret the output as the probability of  $y^{(i)} = 1$



Source: Liang

## Generalizing our gradient descent algorithm

1. Decide a “learning rate”  $\alpha$
2. Start with some parameters  $\theta$
3. For a certain number of iterations
  - Compute  $J(\theta)$
  - Compute  $\nabla_{\theta}J(\theta)$
  - Let  $\theta := \theta - \alpha \nabla_{\theta}J(\theta)$

(initialization)

(forward propagation)

(back-propagation)


(parameter update)




**BAYES**  
BUSINESS SCHOOL  
CITY, UNIVERSITY OF LONDON



## Step 1: Forward propagation through the computational graph



## Step 2: Back-propagation through the computational graph





**Deep learning programming frameworks**

## Why to use programming frameworks for deep learning?

- Build complex neural networks based on simple modules
- Automatically compute gradients
- Make efficient use of hardware (parallel processing, usage of GPUs and TPUs)
- Create models that are exportable to other runtimes, browsers, or even mobile devices
- Load a large range of pre-trained models



# Deep learning in Python: the two main rivals

## TensorFlow

- Developed by Google
- Python API based on C engine
- Largest community
- Parallelization of models and data
- Visualization (TensorBoard)
- Relatively intuitive when familiar with numpy
- Runs on most systems (e.g., TensorFlow Lite)
- Broader framework (e.g., reinforcement learning)
- Easy to use Colab-GPU

## PyTorch

- Developed by Facebook
- Python API based on LUA engine (Torch)
- Many pre-trained models, especially transformers
- Highly modular
- Easy to run on GPUs
- Runs on most systems (e.g., PyTorch Mobile)


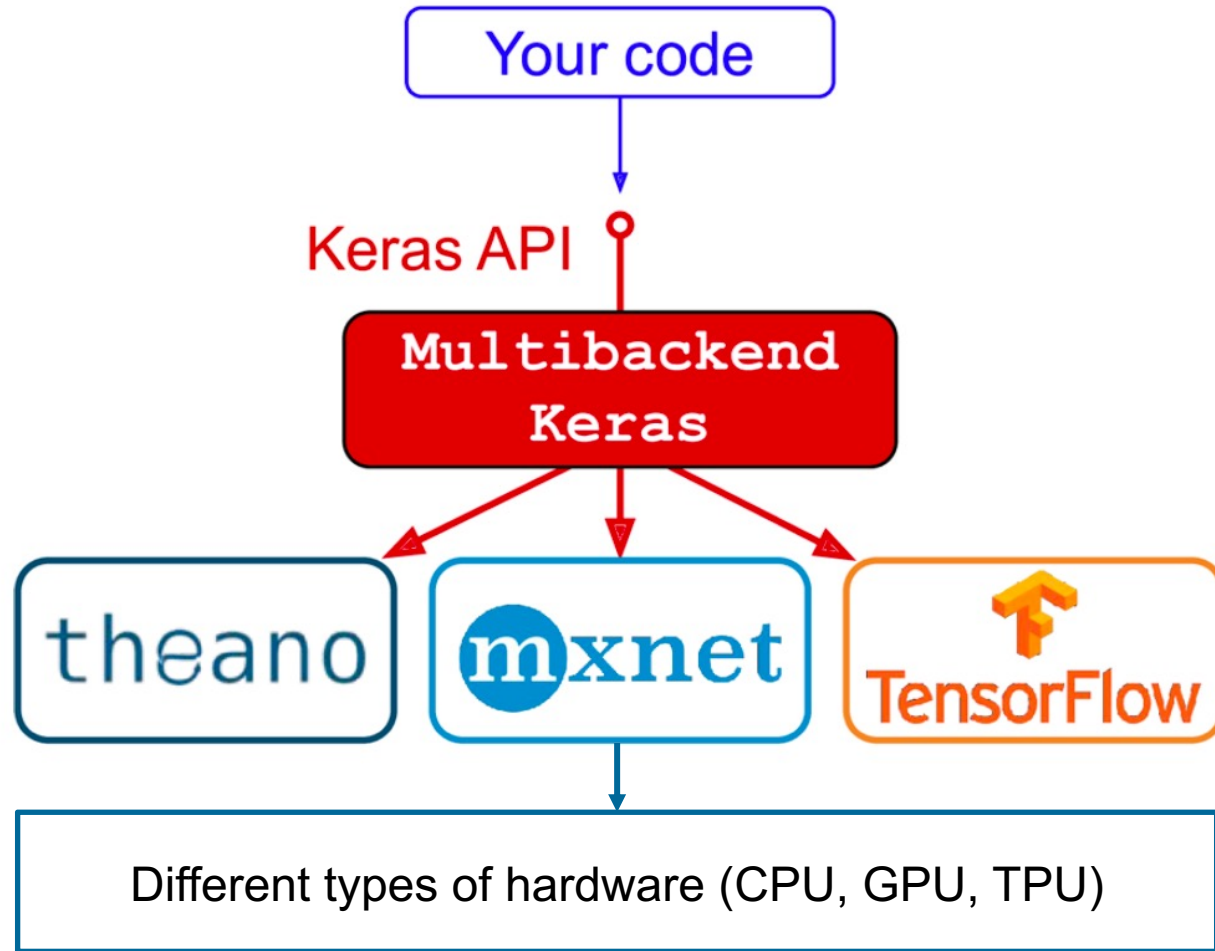


## Wasn't there also this thing called Keras?

- High-level deep learning API to build, train, evaluate, execute neural networks
  - Documentation: <https://keras.io>
- Requires computation backend:
  - TensorFlow (at this point, that's the standard)
  - Microsoft Cognitive Toolkit
  - Theano (original backend on which Keras was built)
  - Apache MXNEt
  - Core ML (Apple)
  - JavaScript & TypeScrit (for web browsers)
  - PlaidML
  - ...
- TensorFlow comes with its own Keras implementation tf.keras (with some added features)
  - E.g., use Data API from TF



## Keras in general and Keras using TensorFlow




Source: Geron




**BAYES**  
BUSINESS SCHOOL  
CITY UNIVERSITY OF LONDON

## Some statistics

ML tools used by top-5 teams in Kaggle competitions 2017-2019



ML tools used by professional data scientists surveyed by Kaggle



Source: Kaggle



**BAYES**  
BUSINESS SCHOOL  
CITY, UNIVERSITY OF LONDON




# TensorFlow in practice



**BAYES**  
BUSINESS SCHOOL  
CITY, UNIVERSITY OF LONDON

## Multi-class classification


- We again make a linear transformation with a matrix of weights:  $\mathbf{z}^{[L](i)} = \mathbf{a}^{[L-1](i)} \mathbf{W}^{[L]} + \mathbf{b}^{[L]}$
- Note that  $\mathbf{z}^{[L](i)} = (z_1^{[L](i)} \quad z_2^{[L](i)} \quad \dots \quad z_K^{[L](i)})$
- We then use the softmax function on each of the outputs:
$$\hat{y}_k^{(i)} = f(\mathbf{z}^{[L](i)}) = \frac{e^{-z_k^{[L](i)}}}{\sum_{k=1}^K e^{-z_k^{[L](i)}}}$$
- This implies that  $\hat{y}_k^{(i)} \in (0,1)$  and  $\sum_{k=1}^K \hat{y}_k^{(i)} = 1$
- Hence, we can interpret  $\hat{y}_k^{(i)}$  as the probability that  $y^{(i)} = k$  (“belongs to class  $k$ ”)



Source: Liang

## Softmax output

- E.g., when performing object recognition, we might represent our prediction  $\hat{y}^{(i)}$  as



## Softmax in practice



**BAYES**  
BUSINESS SCHOOL  
CITY, UNIVERSITY OF LONDON



See you next week!

## Sources

- Chollet, 2021, Deep Learning with Python (2<sup>nd</sup> edition)
- DeepLearning.AI, n.d.: [deeplearning.ai](https://www.deeplearning.ai)
- Geron, 2019, Hands-on Machine Learning with Scikit-Learn, Keras, and TensorFlow
- Goodfellow, Bengio, Courville, 2016, The Deep Learning Book:  
<http://www.deeplearningbook.org>
- Kaggle, 2019, Kaggle competitions: Why choose Keras as framework?:  
<https://www.kaggle.com/getting-started/151903>
- Kaggle, 2022, State of Data Science and Machine Learning 2022:  
<https://www.kaggle.com/kaggle-survey-2022>
- Liang, 2016, Introduction to Deep Learning:  
<https://www.cs.princeton.edu/courses/archive/spring16/cos495/>

