



# Applied Deep Learning

Dr. Philippe Blaettchen  
Bayes Business School (formerly Cass)

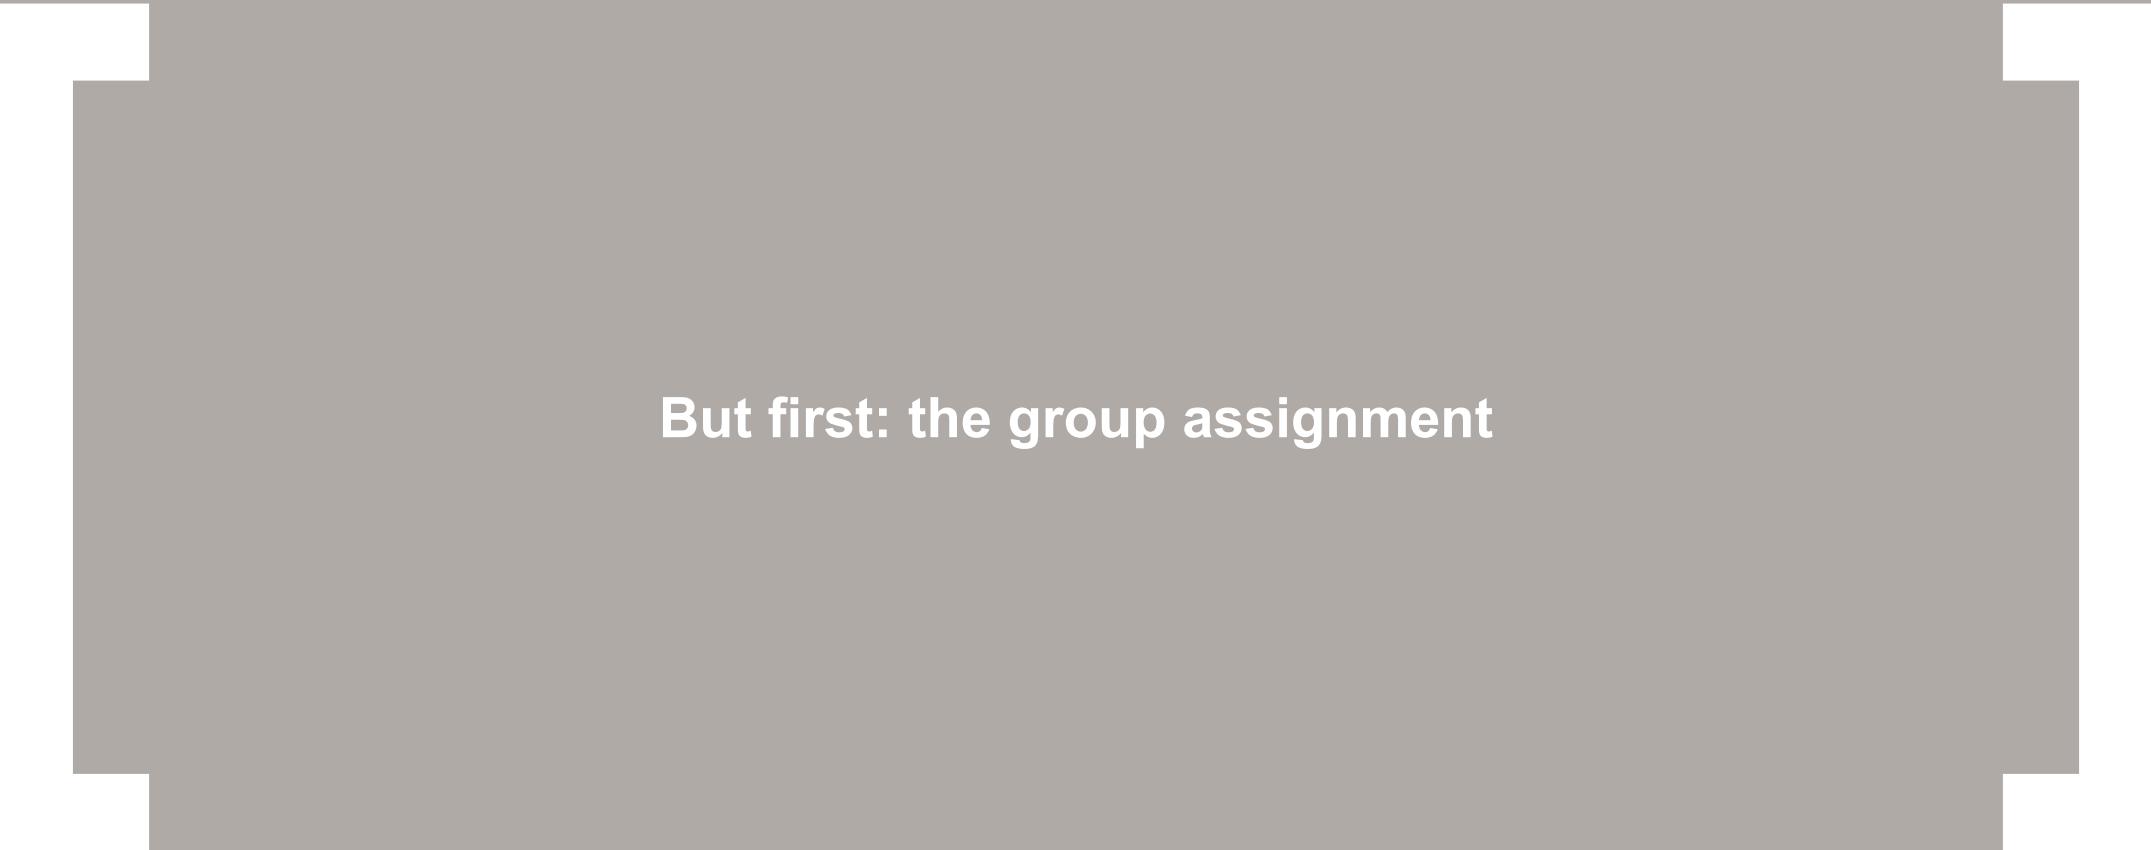
# Learning objectives of today

## Goals:

- Discuss the importance of computer vision for different business applications
- Understand how convolution works and how it can be implemented in TensorFlow
- Apply CNNs to medical diagnosis

## How will we do this?

- We will start with a look at Zebra Medical Vision, which provides diagnostic services based on medical imaging
- We will see how convolutional layers drastically simplify the task of analyzing images
- We then turn to the impact of developments in computer vision more broadly



**But first: the group assignment**

# The importance of transparency

## The importance of transparency



Trust

## The importance of transparency



Trust



Fairness

## The importance of transparency



Trust



Fairness



Regulation

## The importance of transparency



Trust



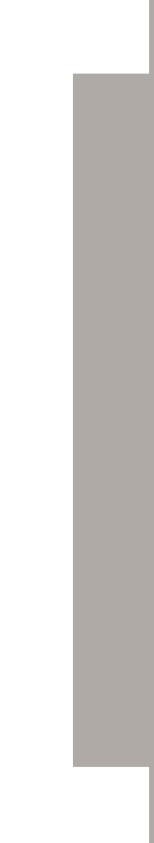
Fairness



Regulation



Informativeness



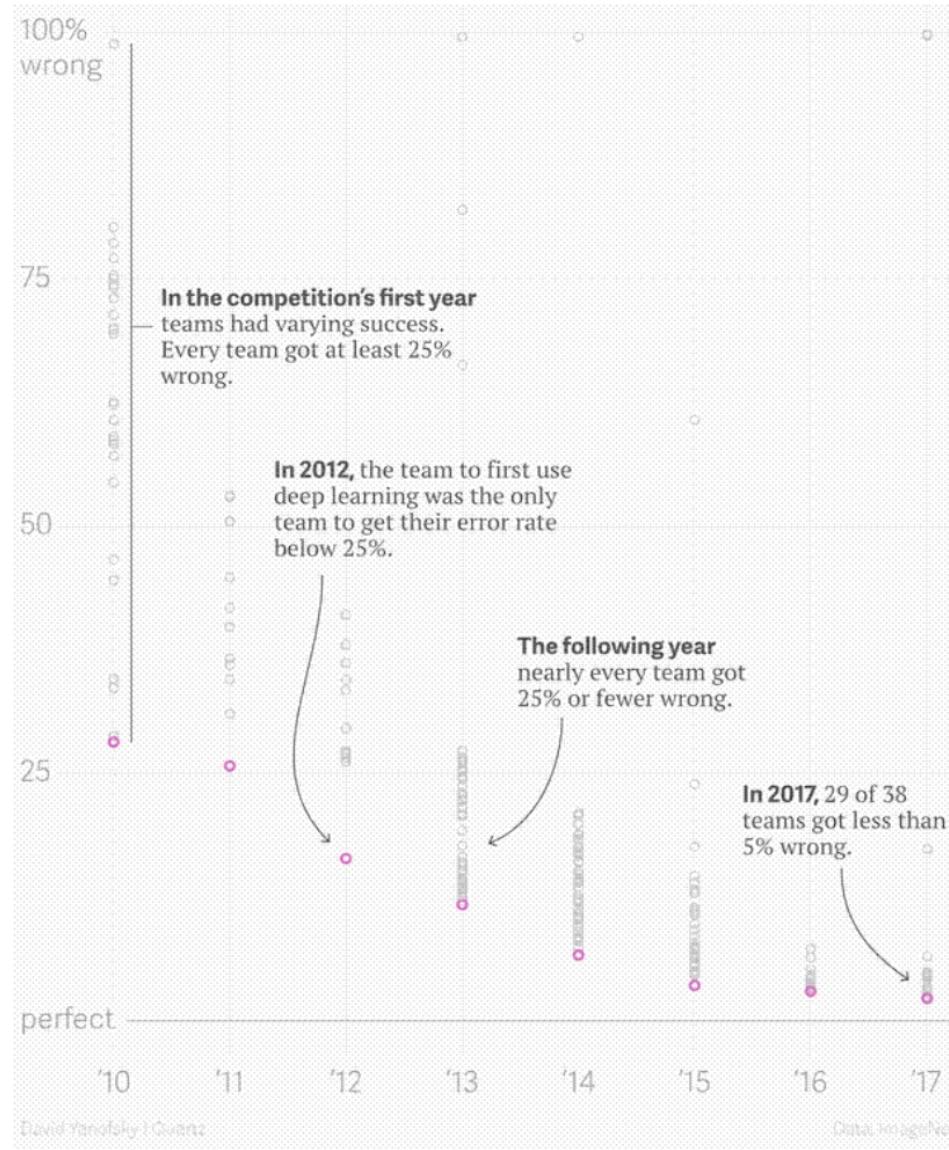
**Zebra Medical Vision**

# The startup

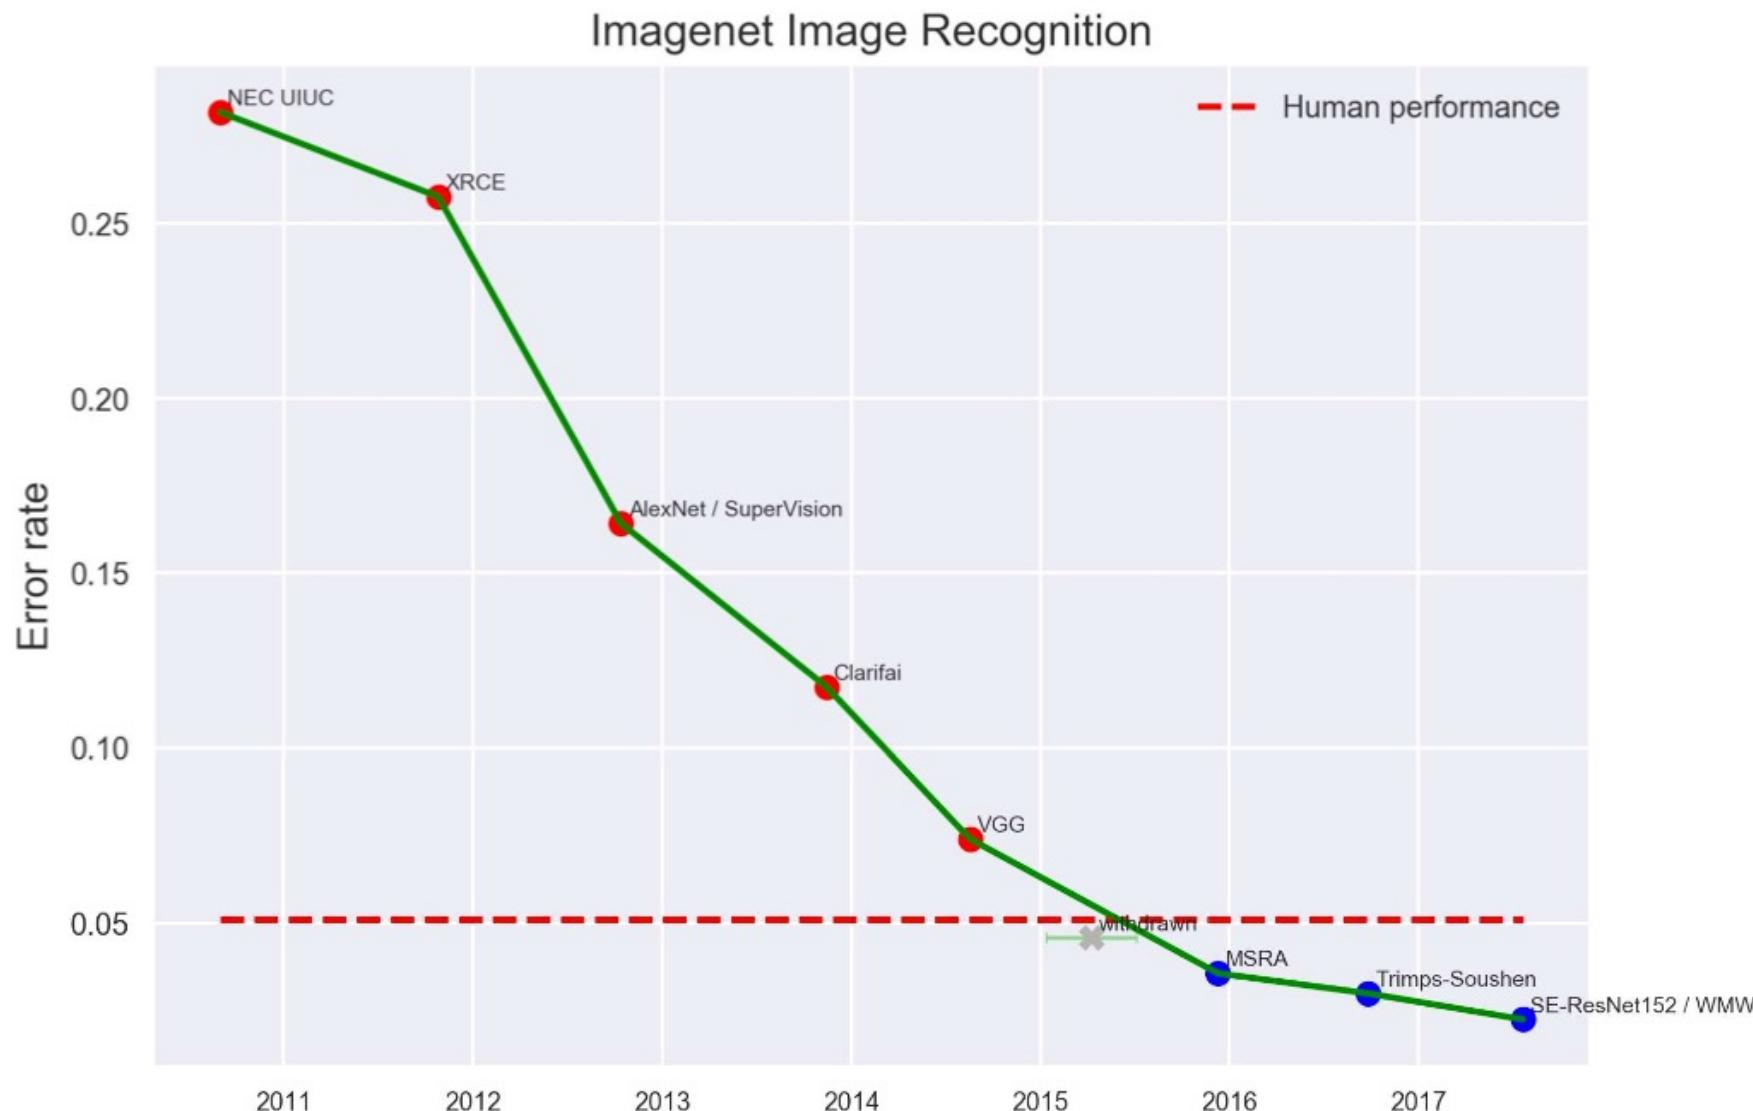
Why now? Why not a decade ago?



# ImageNet Large Scale Visual Recognition Challenge Results



# ImageNet Large Scale Visual Recognition Challenge Results



Why now? Why not a decade ago?

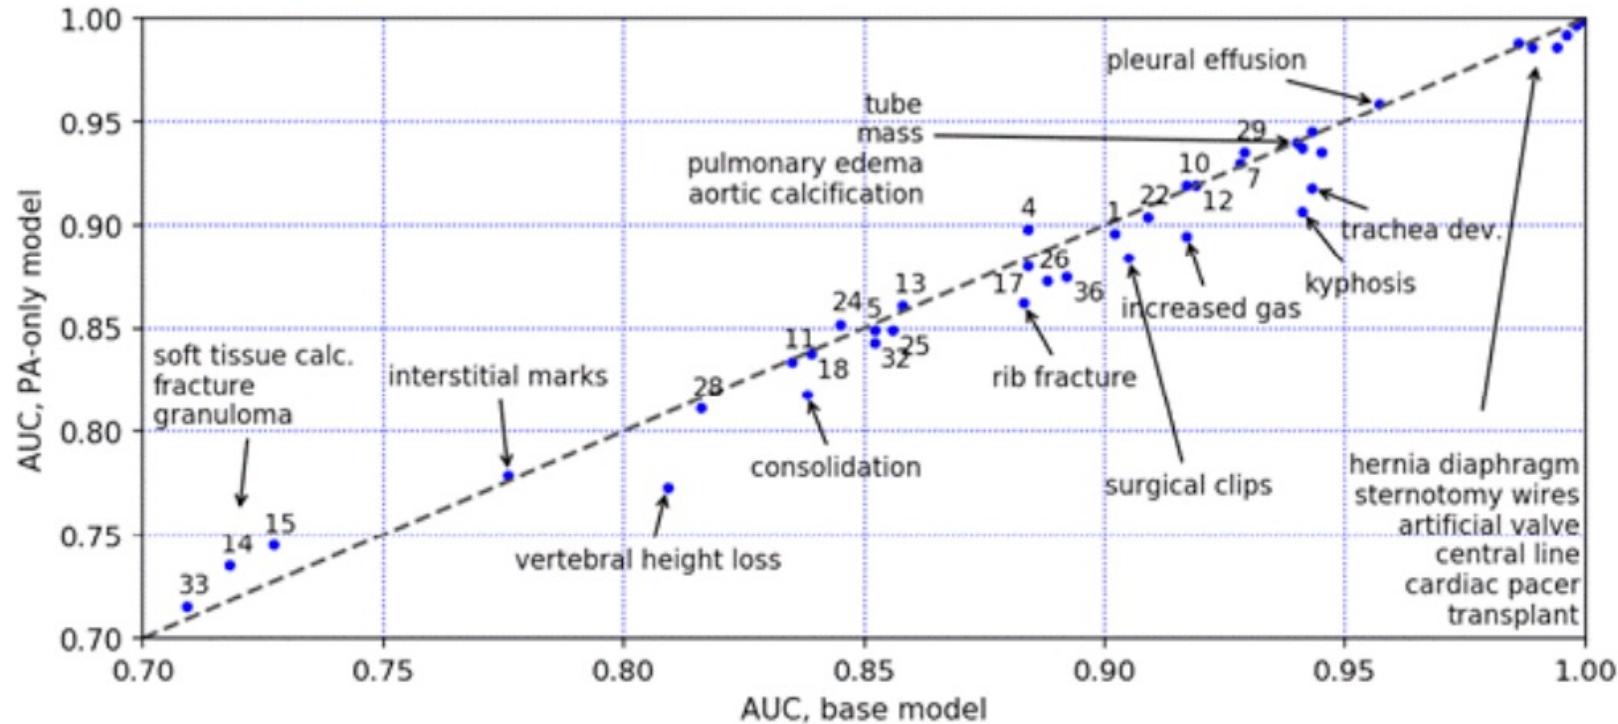


# What is Zebra fundamentally doing?

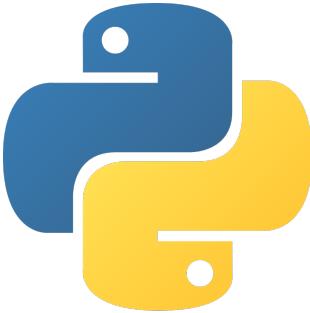
# What is Zebra fundamentally doing?

		True state	
		Positive	Negative
Detected state	Positive	<b>True positive</b>	<b>False positive</b>
	Negative	<b>False negative</b>	<b>True negative</b>

## Predictions based on lung X-rays



Let's try classifying lung X-rays ourselves

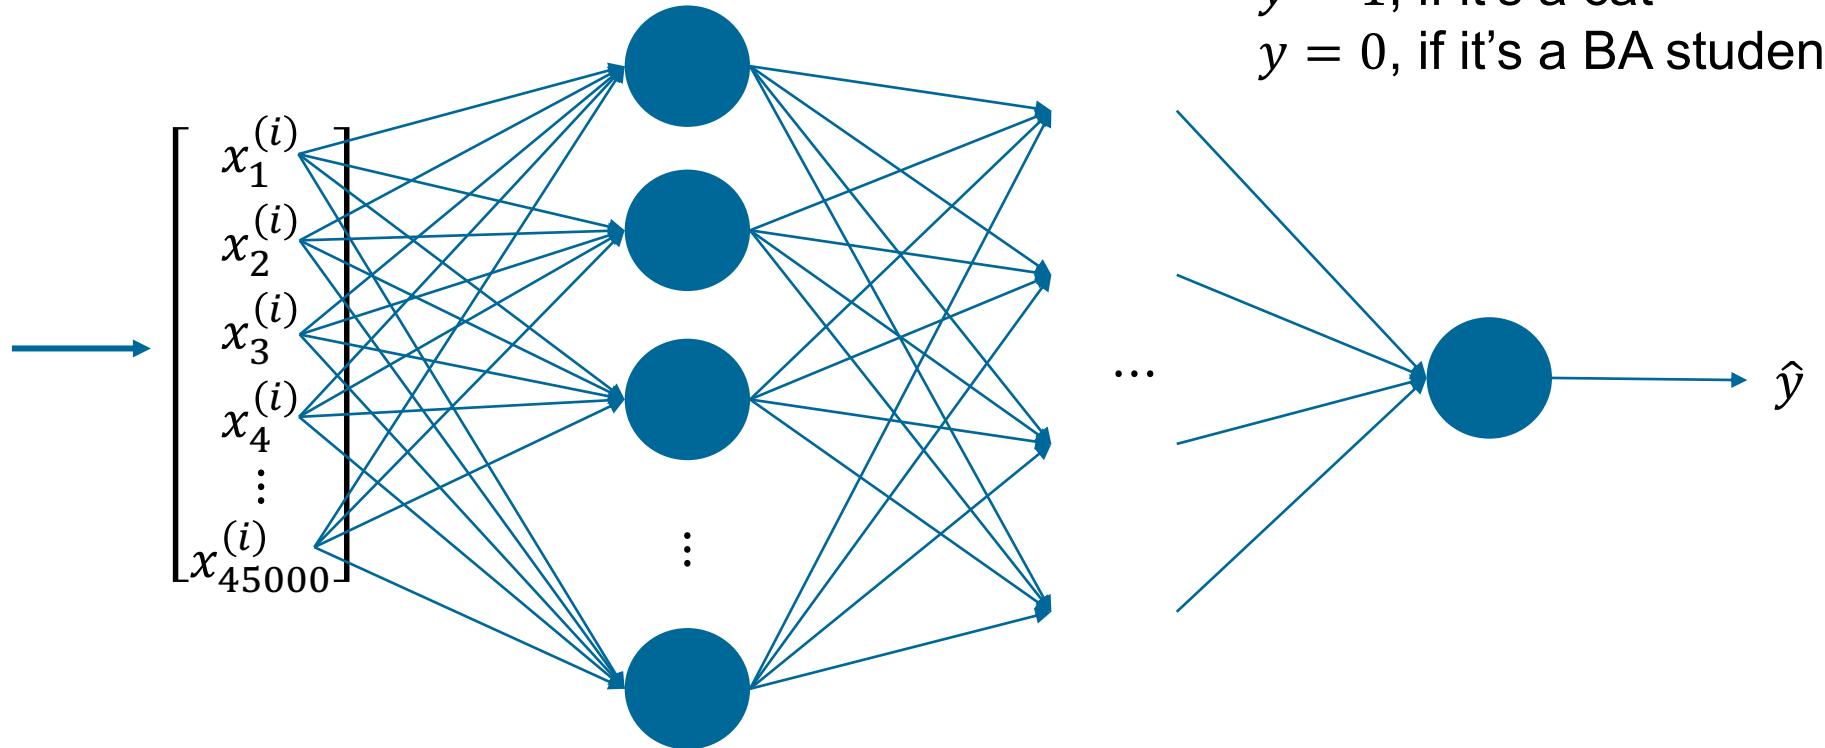


## **Basics of convolution**

## Challenges to deep learning using large images



150x100x3

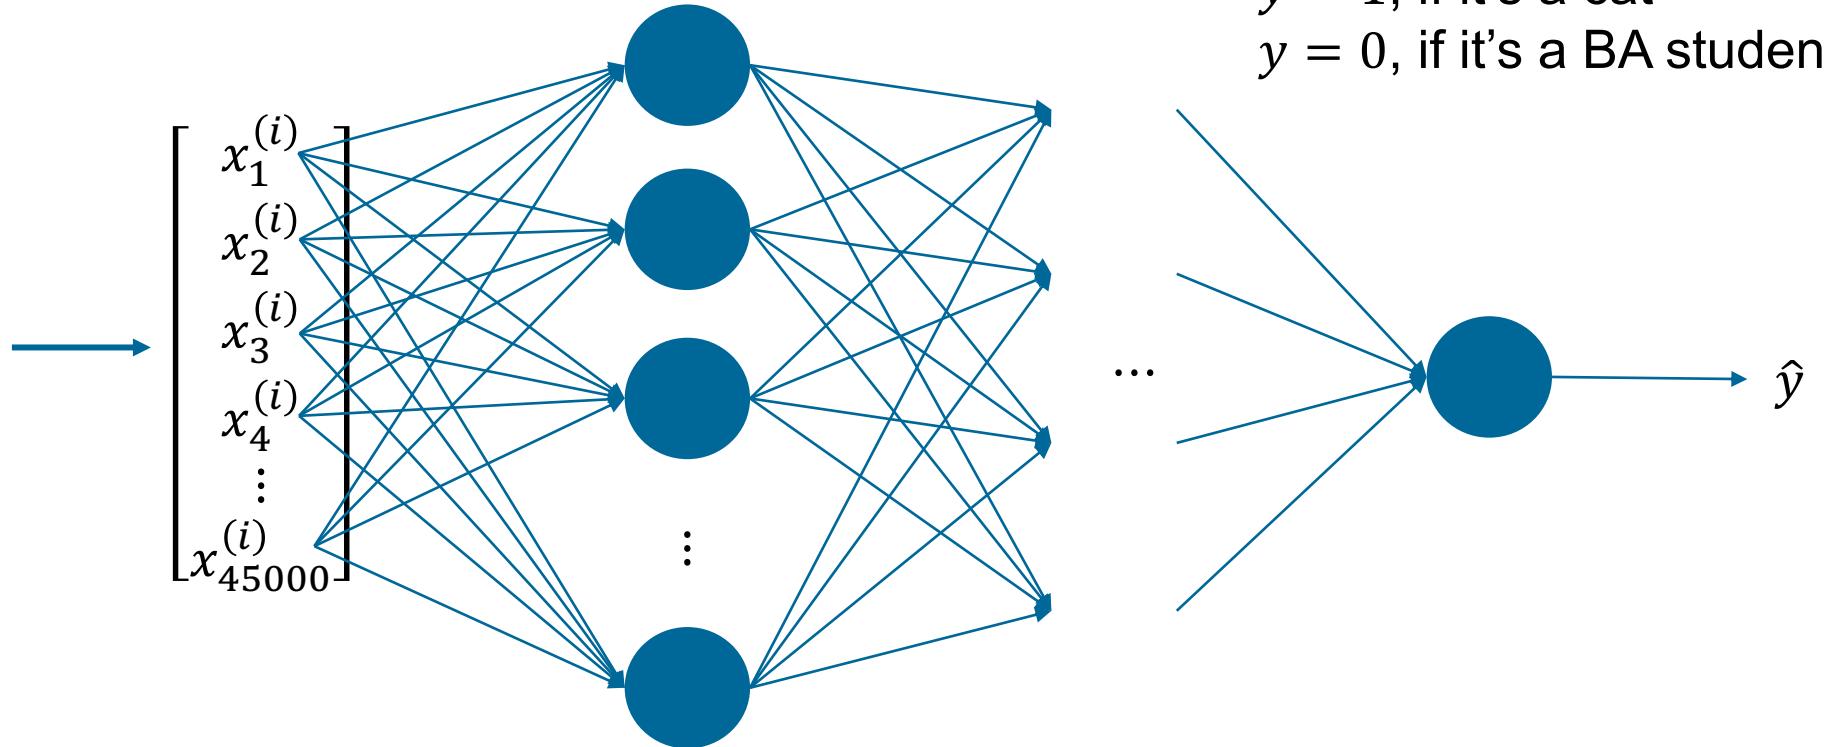


$y = 1$ , if it's a cat  
 $y = 0$ , if it's a BA student

## Challenges to deep learning using large images



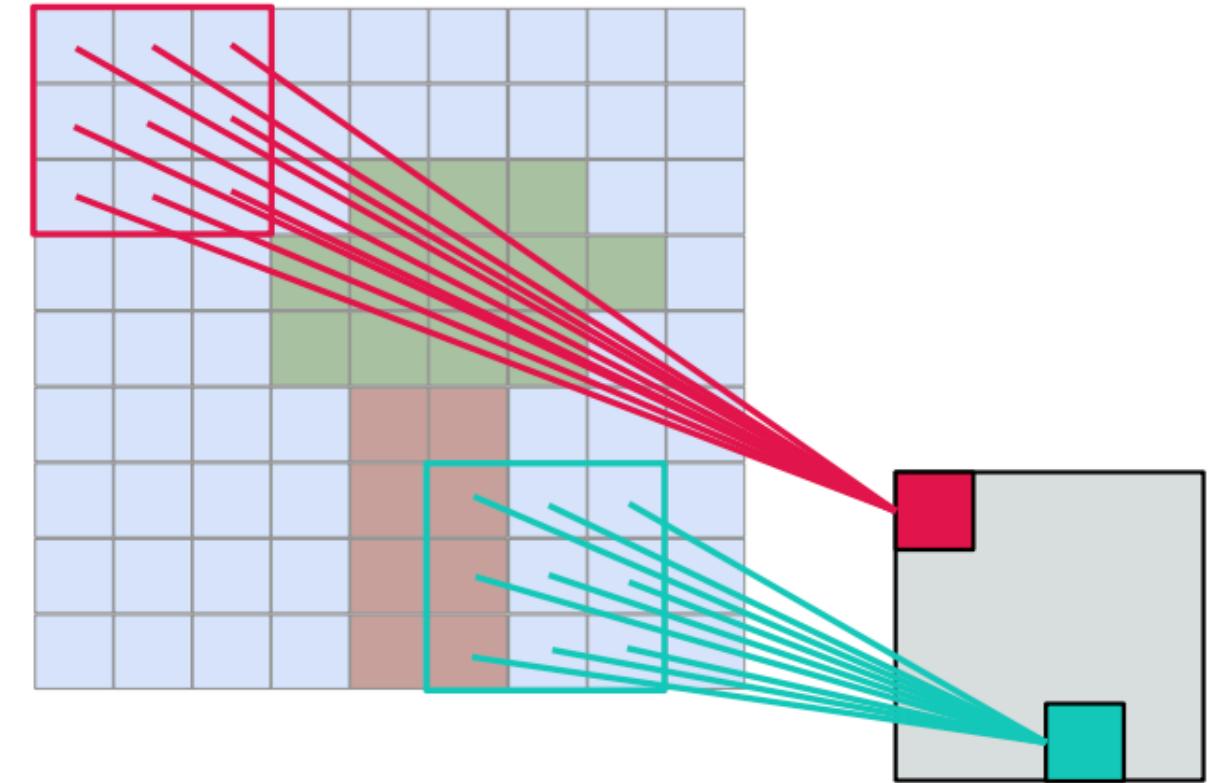
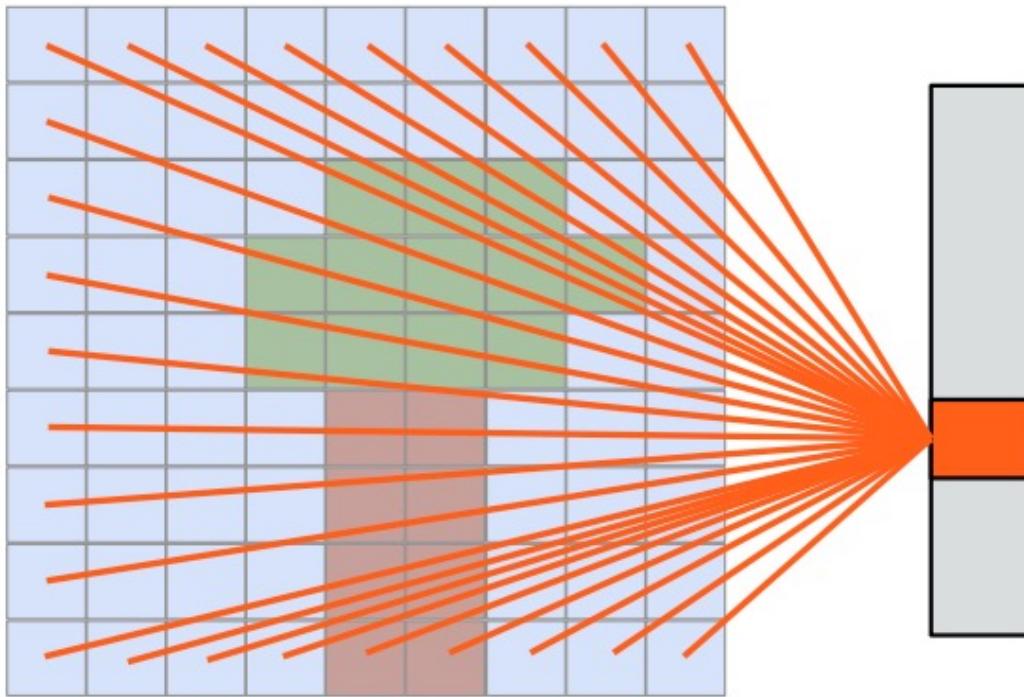
150x100x3



Say we have 100 neurons at the first layer.

Then we need  $100 * (45000 + 1) = 4.5$  mio parameters, just for the first layer!

## From fully connected to locally connected



Source: Dieleman

## An image filter

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

\*

1	0	-1
1	0	-1
1	0	-1

=


## An image filter

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

\*

1	0	-1
1	0	-1
1	0	-1

=

0			

## An image filter

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

\*

1	0	-1
1	0	-1
1	0	-1

=

0			

## An image filter

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

\*

1	0	-1
1	0	-1
1	0	-1

=

0	30		

## An image filter

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

\*

1	0	-1
1	0	-1
1	0	-1

=

0	30		

## An image filter

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

\*

1	0	-1
1	0	-1
1	0	-1

=

0	30	30	

## An image filter

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

\*

1	0	-1
1	0	-1
1	0	-1

=

0	30	30	

## An image filter

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

\*

1	0	-1
1	0	-1
1	0	-1

=

0	30	30	0

## An image filter

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

\*

1	0	-1
1	0	-1
1	0	-1

=

0	30	30	0
0	30	30	0
0	30	30	0

## An image filter

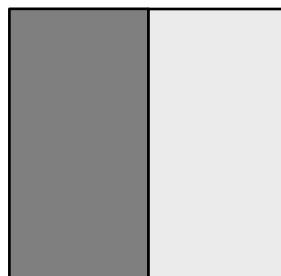
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

\*

1	0	-1
1	0	-1
1	0	-1

=

0	30	30	0
0	30	30	0
0	30	30	0



## An image filter

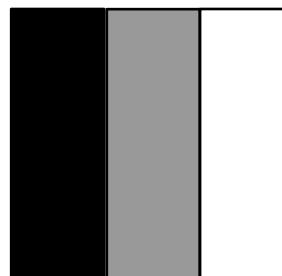
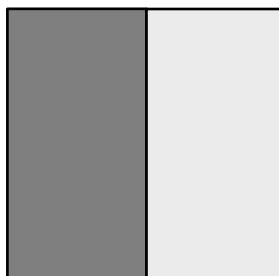
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

\*

1	0	-1
1	0	-1
1	0	-1

=

0	30	30	0
0	30	30	0
0	30	30	0



## An image filter

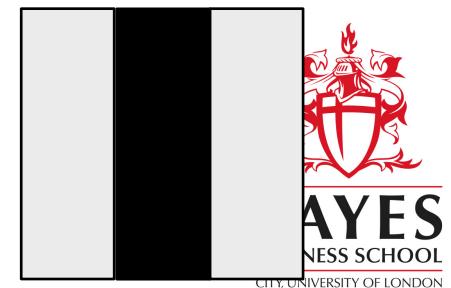
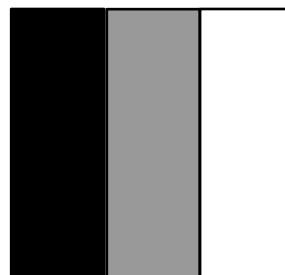
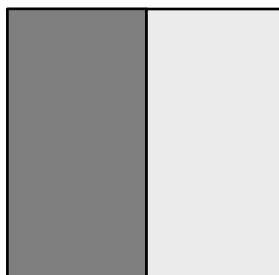
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

\*

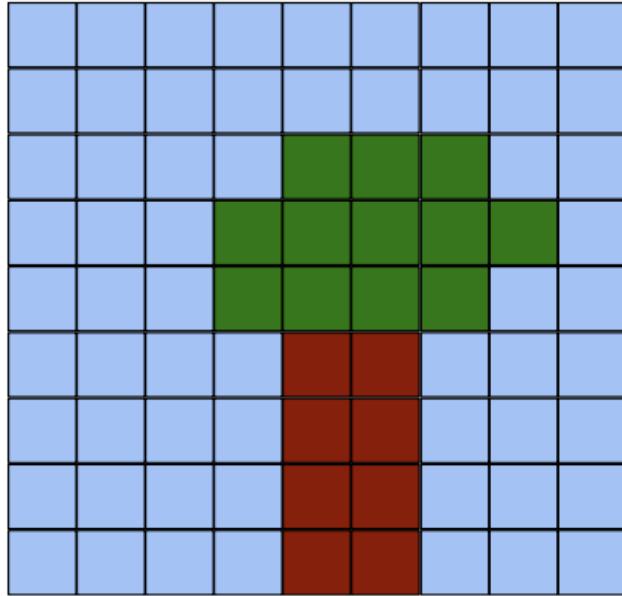
1	0	-1
1	0	-1
1	0	-1

=

0	30	30	0
0	30	30	0
0	30	30	0



## Convolutional layers: learning the filter

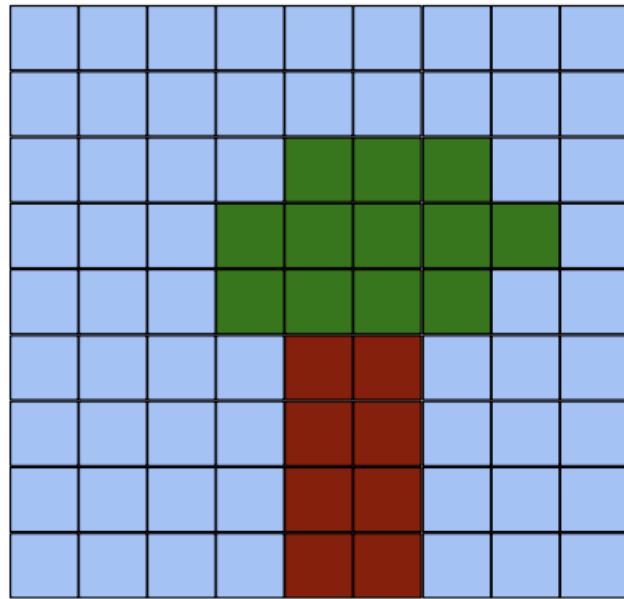


\*

$w_{1,1}$	$w_{1,2}$	$w_{1,3}$
$w_{2,1}$	$w_{2,2}$	$w_{2,3}$
$w_{3,1}$	$w_{3,2}$	$w_{3,3}$

Source: Dieleman

## Convolutional layers: learning the filter



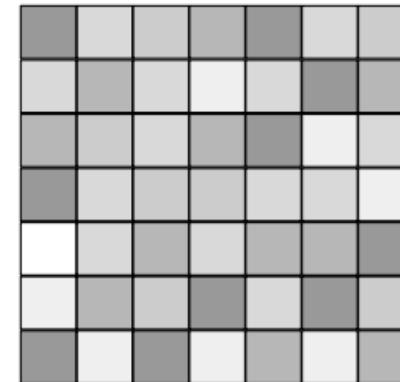
\*

$w_{1,1}$	$w_{1,2}$	$w_{1,3}$
$w_{2,1}$	$w_{2,2}$	$w_{2,3}$
$w_{3,1}$	$w_{3,2}$	$w_{3,3}$

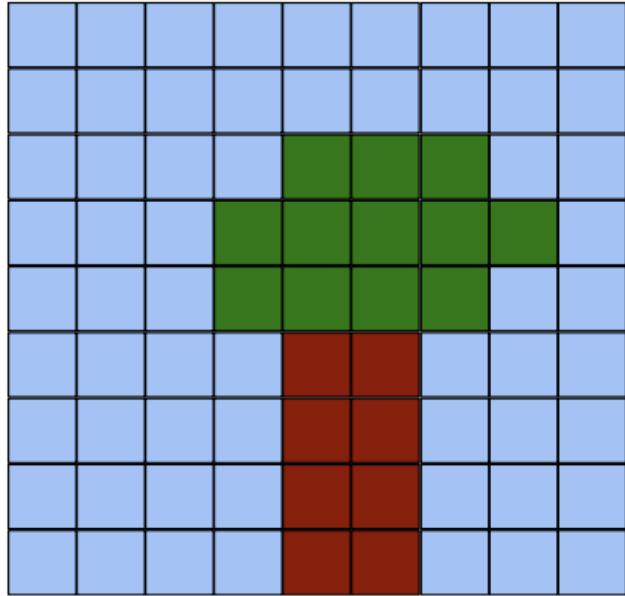
+

$b$

=



## Convolutional layers: learning the filter



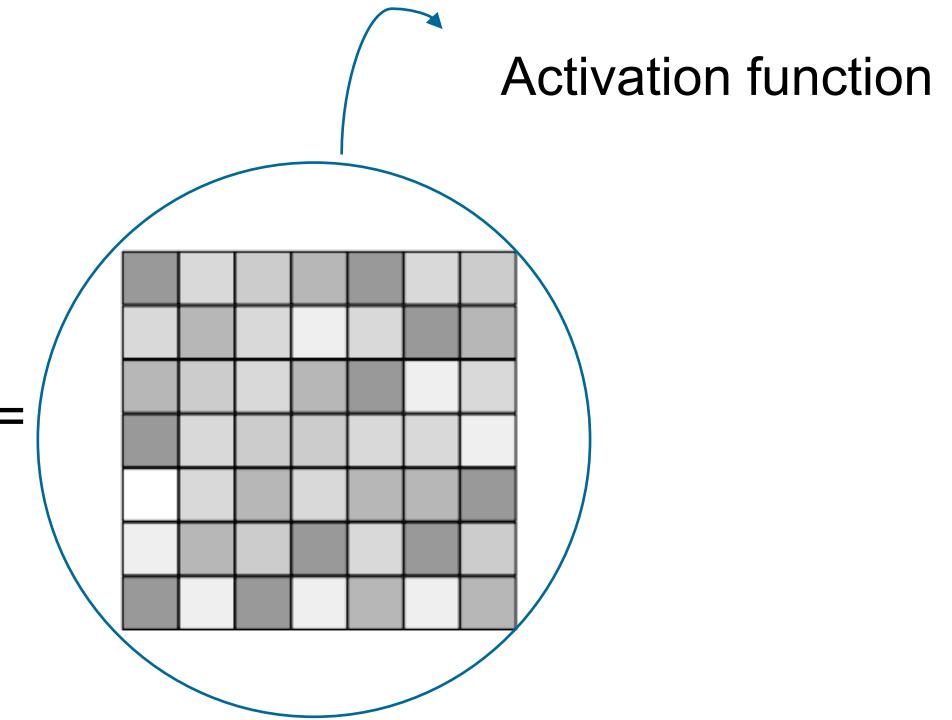
\*

$w_{1,1}$	$w_{1,2}$	$w_{1,3}$
$w_{2,1}$	$w_{2,2}$	$w_{2,3}$
$w_{3,1}$	$w_{3,2}$	$w_{3,3}$

+

$b$

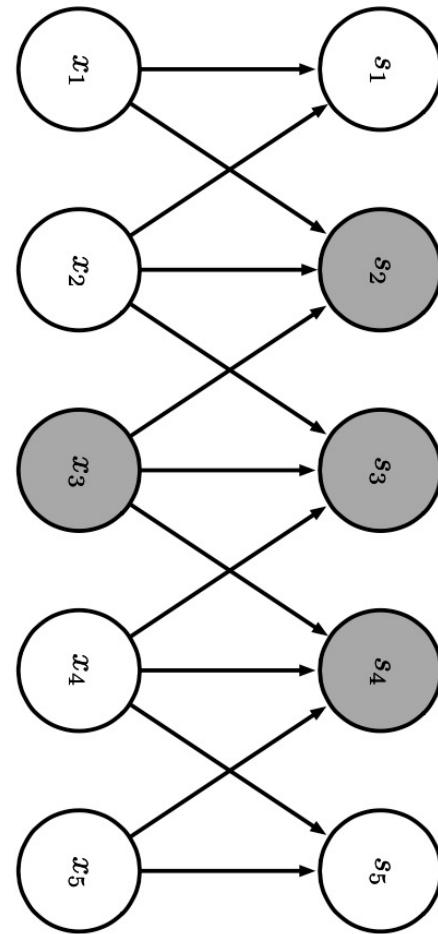
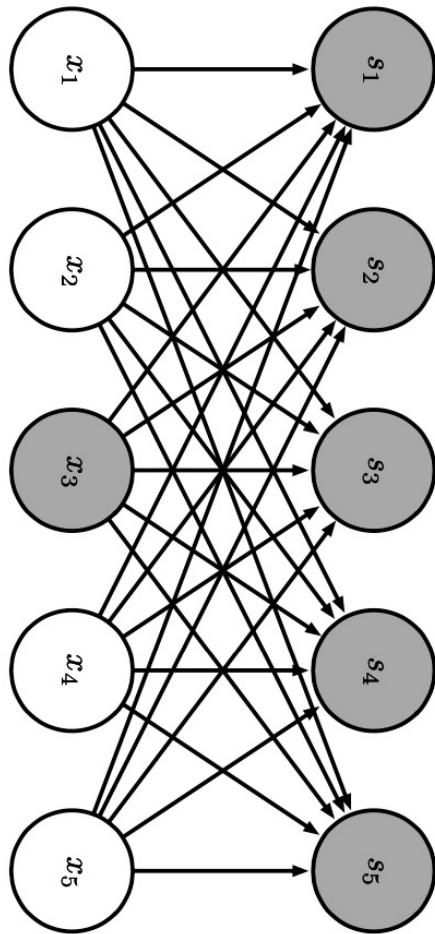
=



Activation function

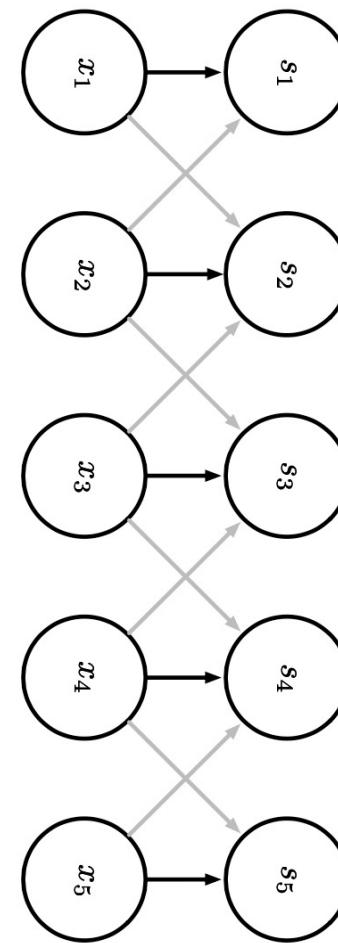
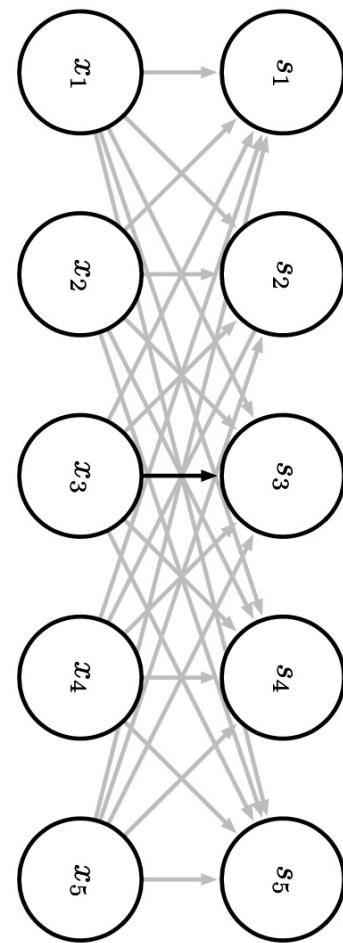
Source: Dieleman

## Less parameters required



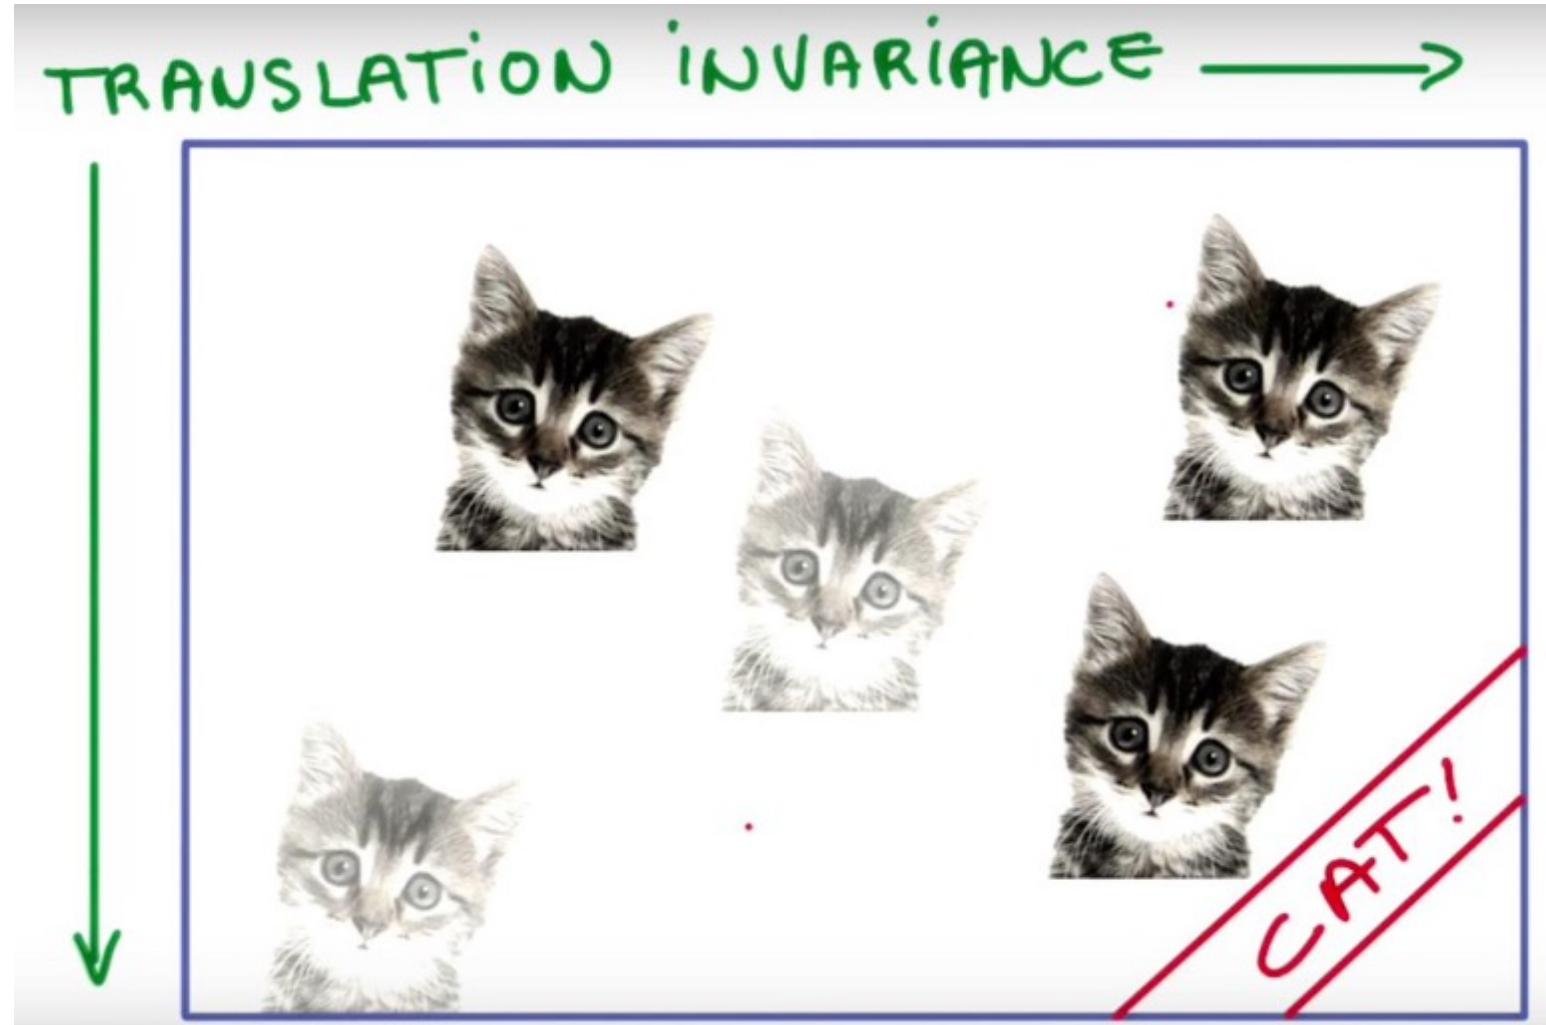
Source: Goodfellow

Parameters are shared



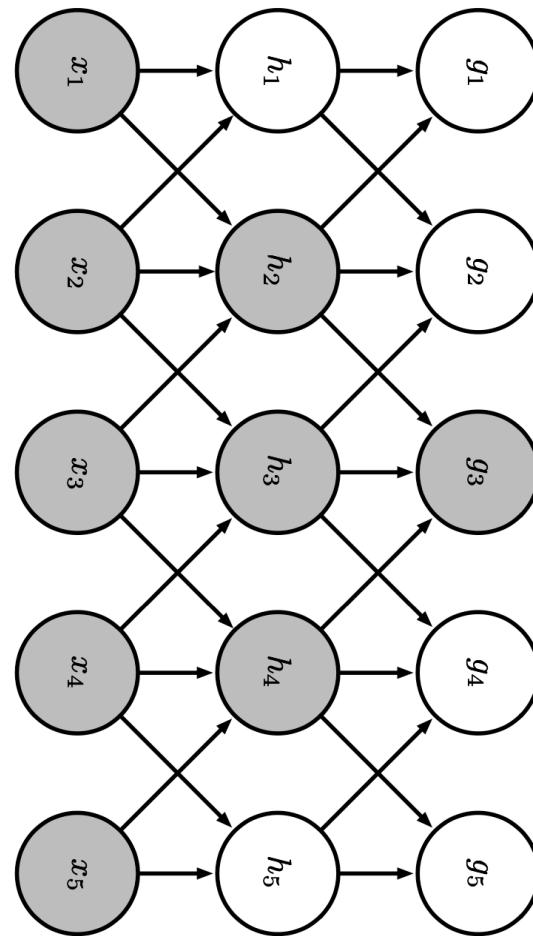
Source: Goodfellow

## Shared parameters lead to translation invariance



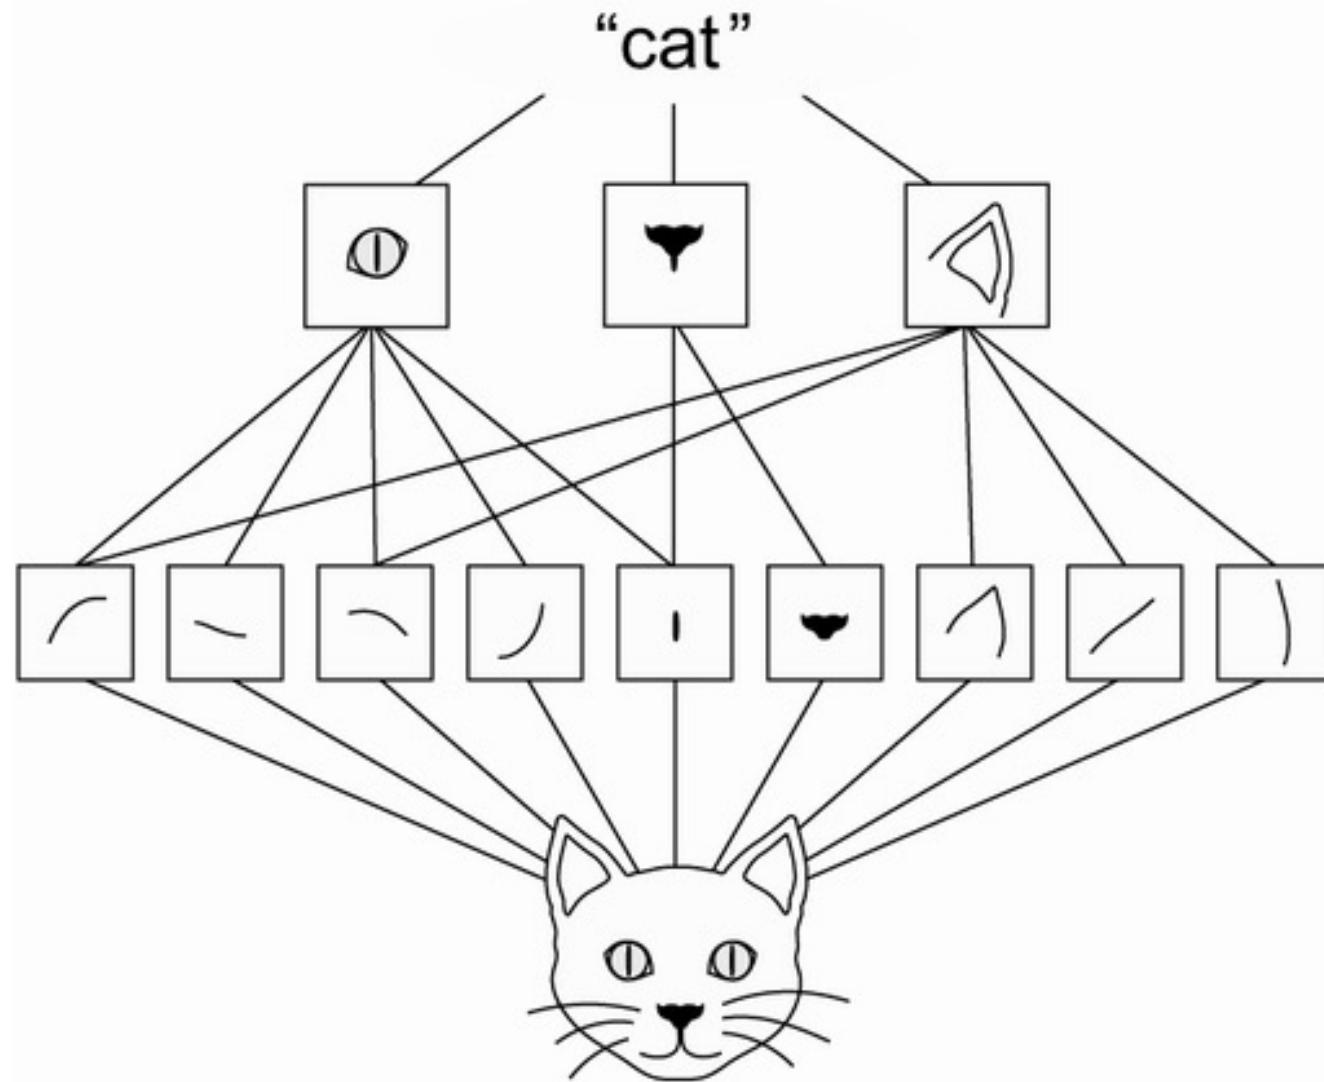
Source: Bhaskhar

## Hierarchical setup



Source: Goodfellow

## Hierarchical structure supports a spatial hierarchy of learning

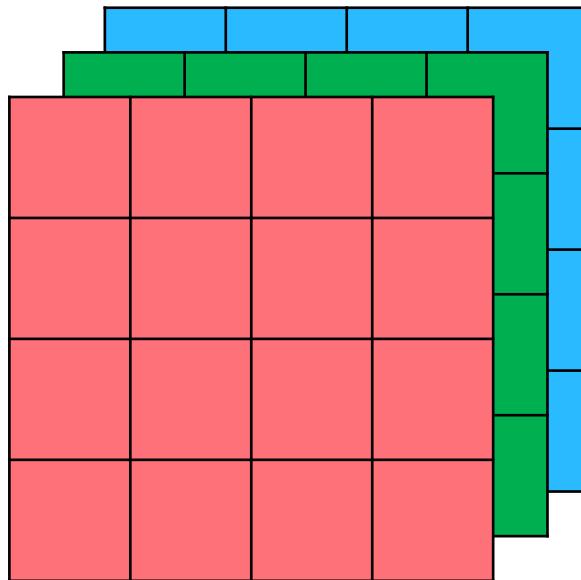


Source: Chollet

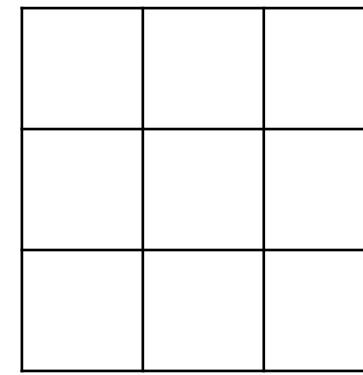


**From convolution to CNNs**

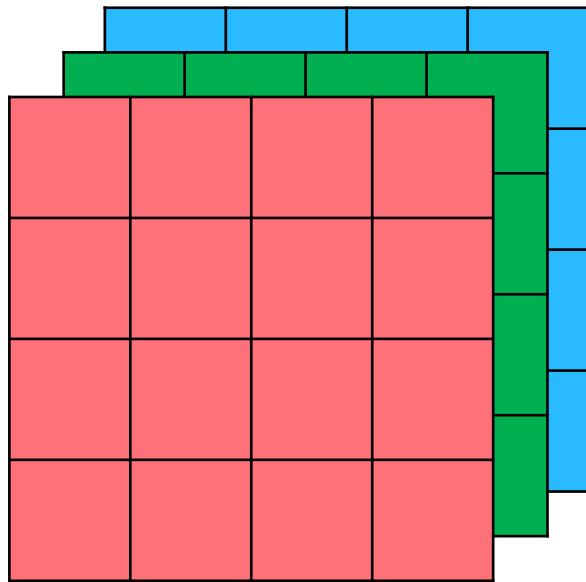
## Convolution on a 3D array



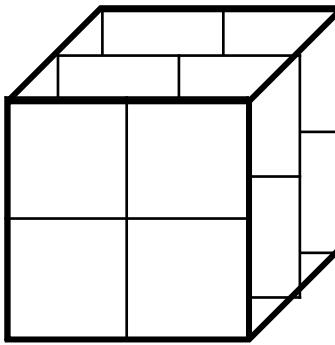
=



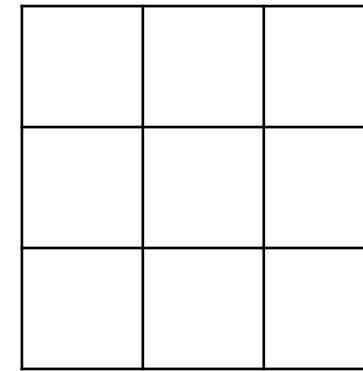
## Convolution on a 3D array



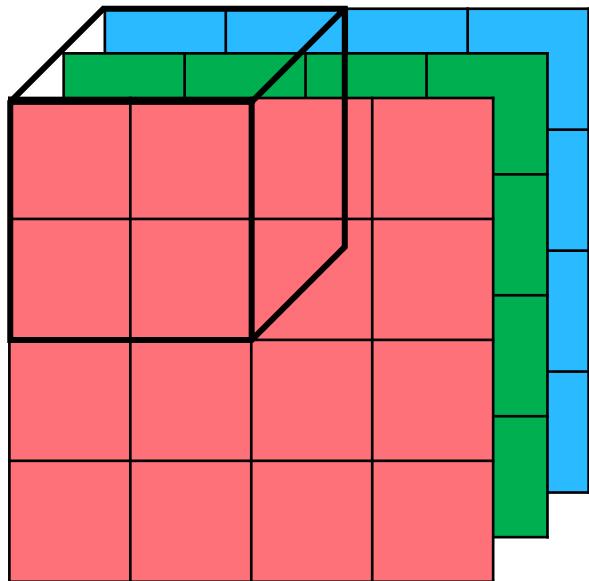
\*



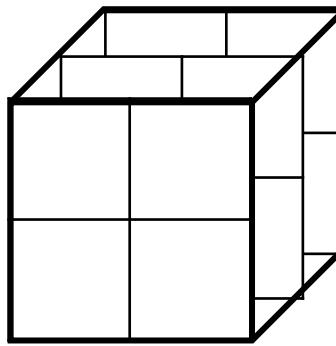
=



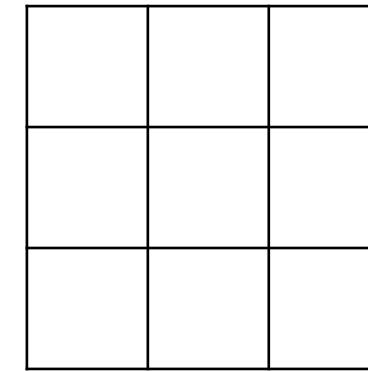
## Convolution on a 3D array



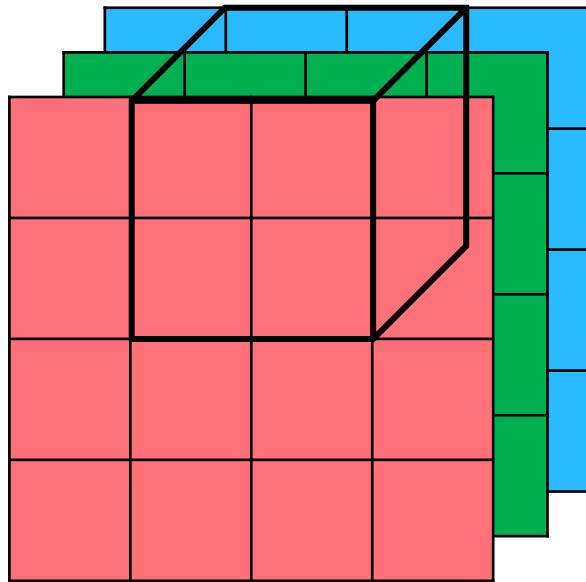
\*



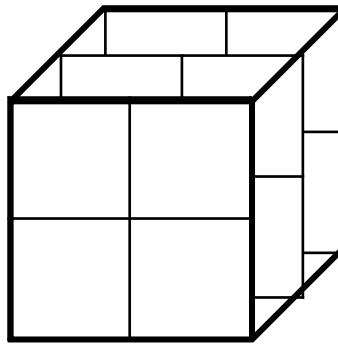
=



## Convolution on a 3D array



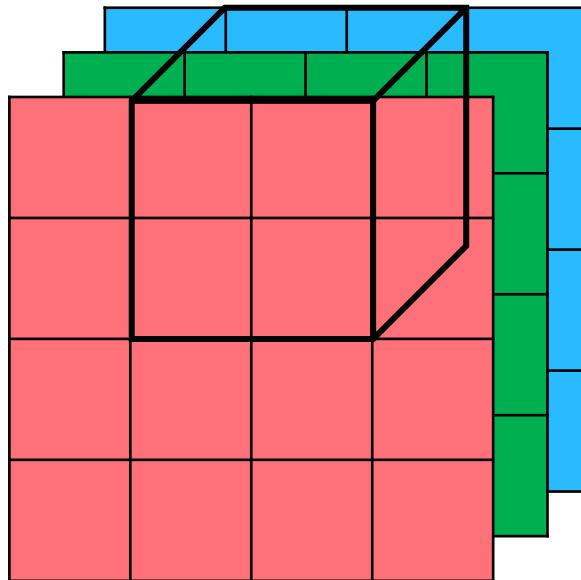
\*



=

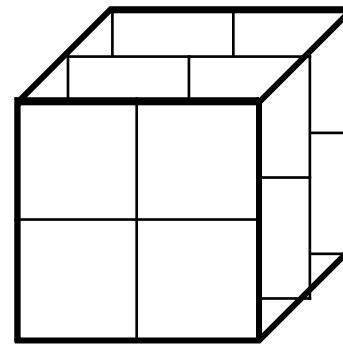
20		

## Convolution on a 3D array



4x4x3

\*



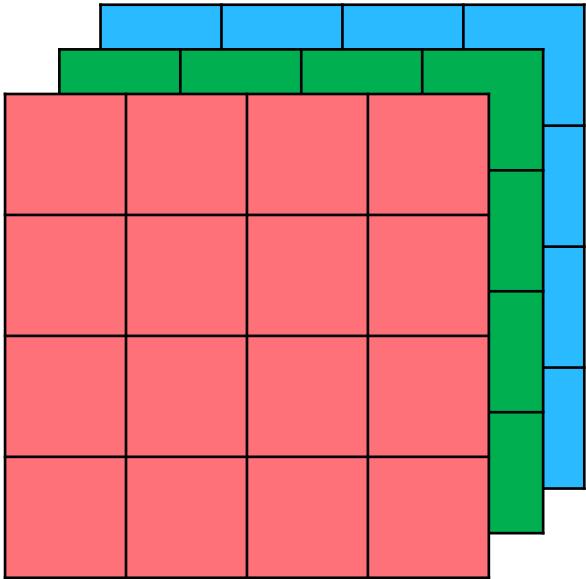
2x2x3

=

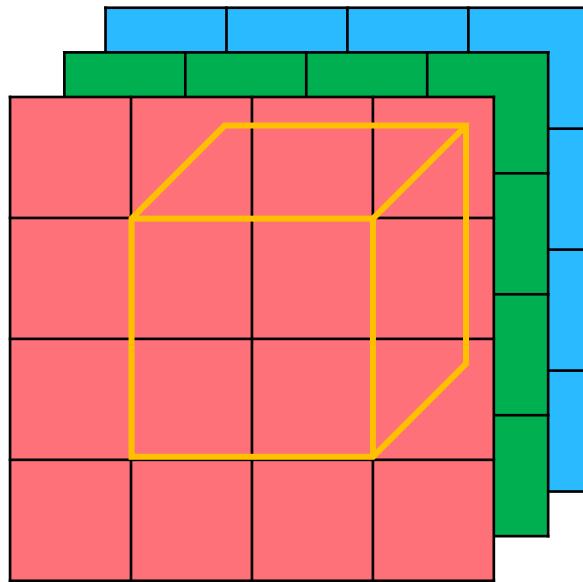
20		

3x3

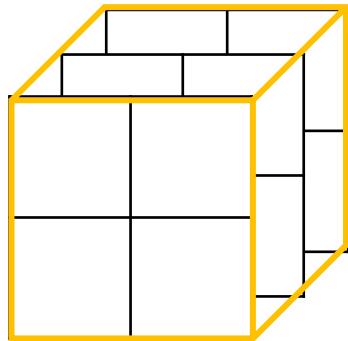
# Multiple 3D convolutions



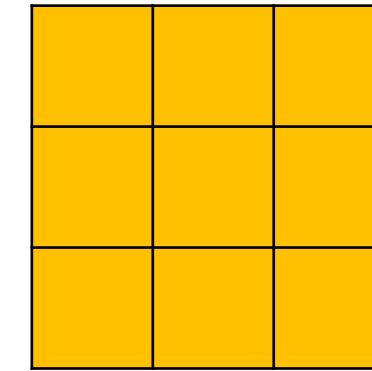
## Multiple 3D convolutions



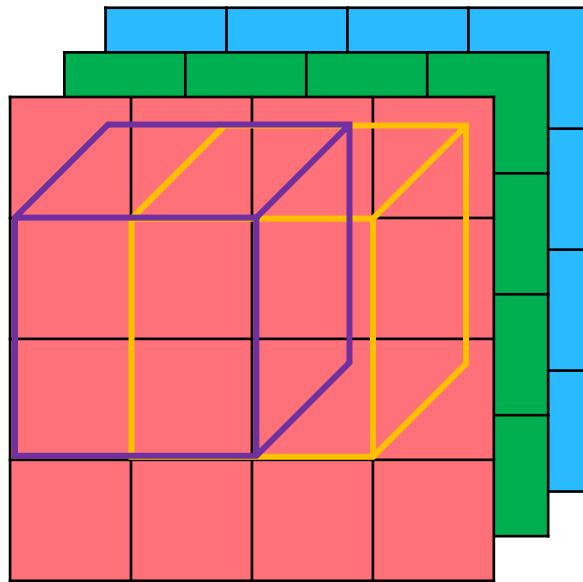
\*



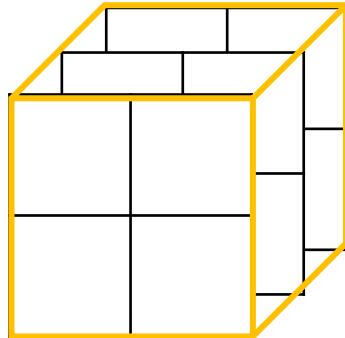
=



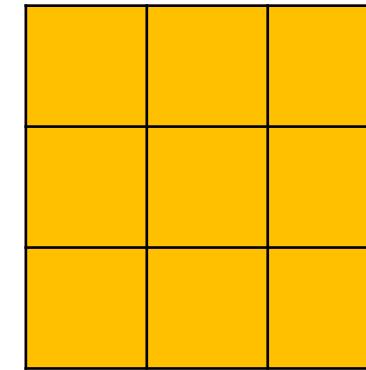
## Multiple 3D convolutions



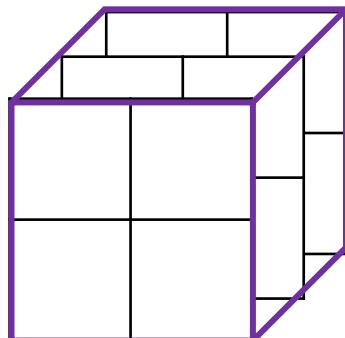
\*



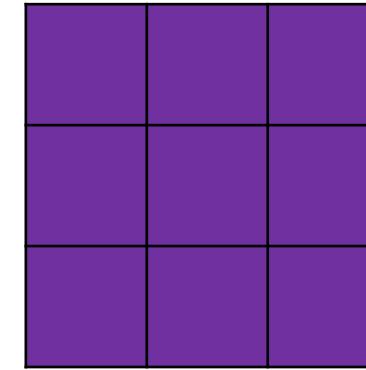
=



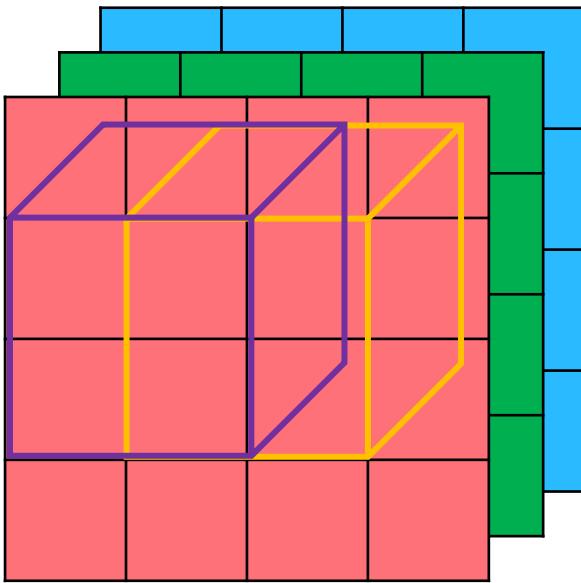
\*



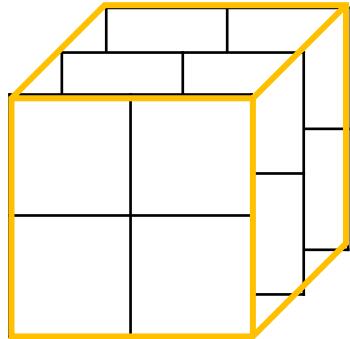
=



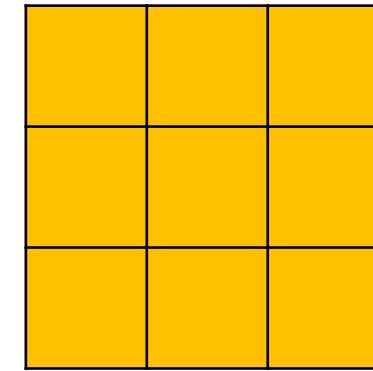
## Multiple 3D convolutions



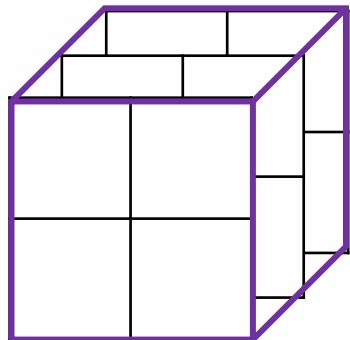
\*



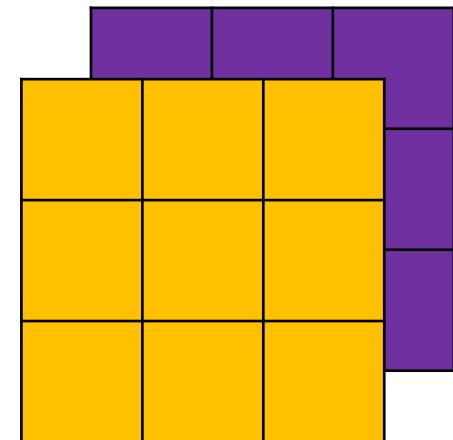
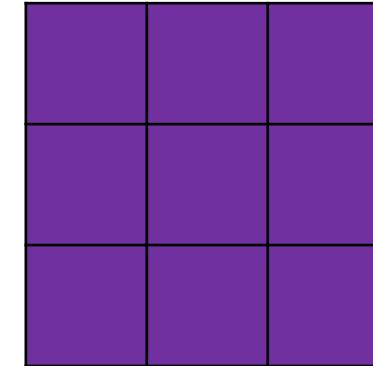
=



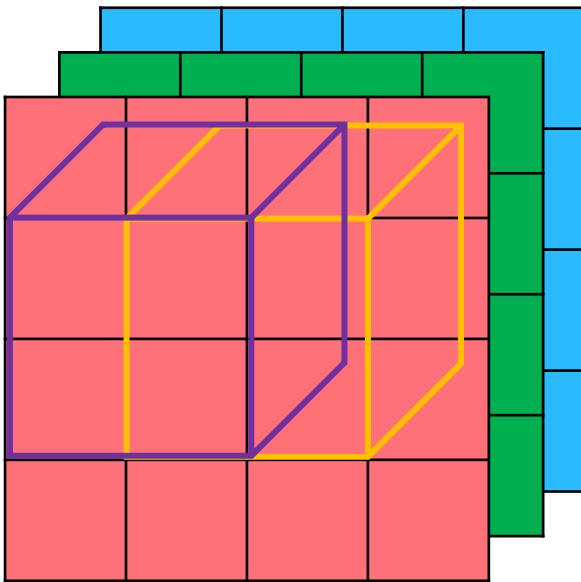
\*



=

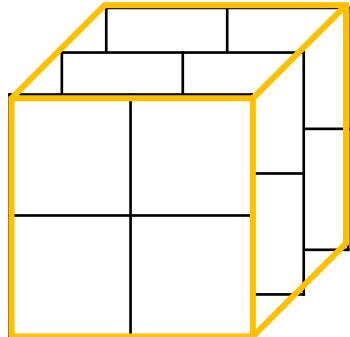


## Multiple 3D convolutions

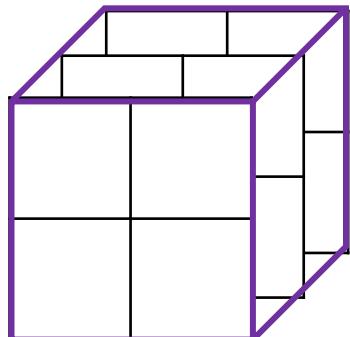
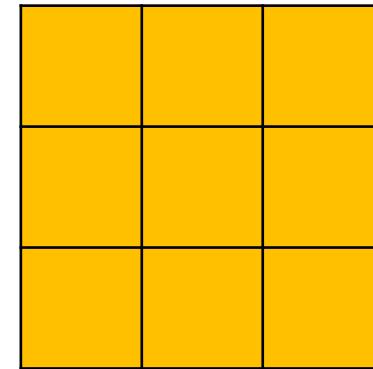


$4 \times 4 \times 3$   
 $(n_H^{[0]}, n_W^{[0]}, n_C^{[0]})$

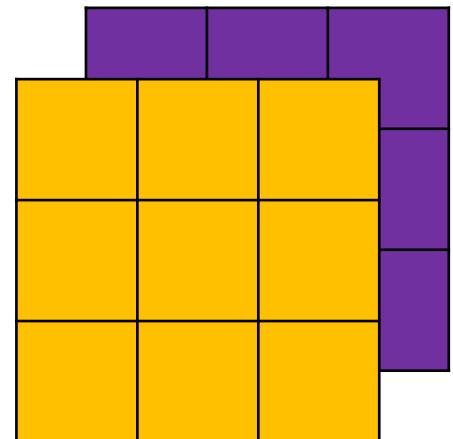
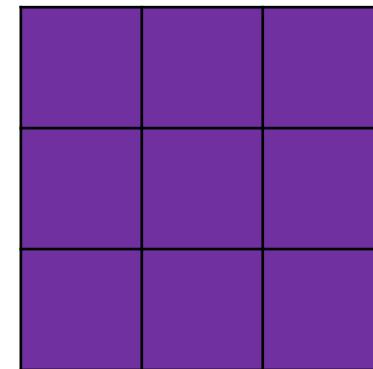
\*



=

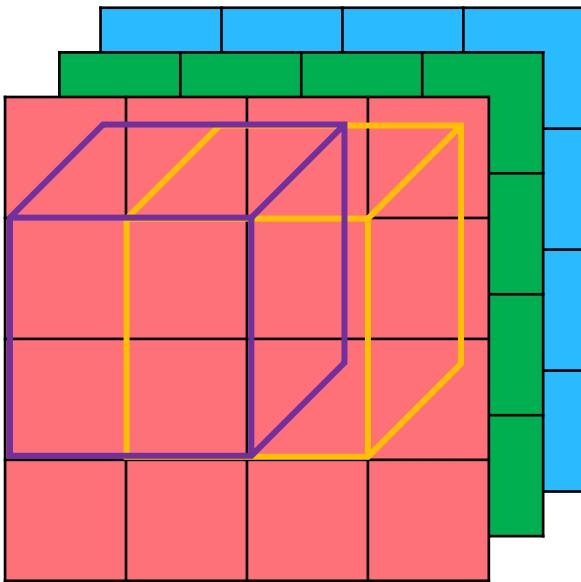


=



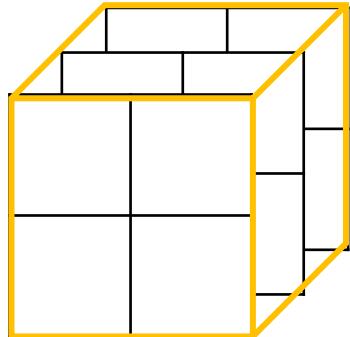
$n_C^{[1]} = \text{number of filters}$

## Multiple 3D convolutions

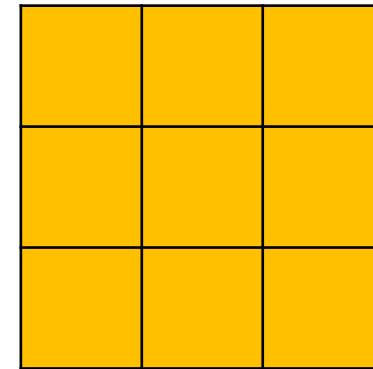


$4 \times 4 \times 3$   
 $(n_H^{[0]}, n_W^{[0]}, n_C^{[0]})$

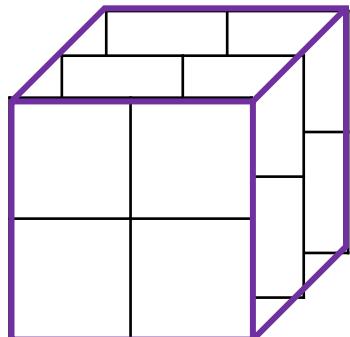
\*



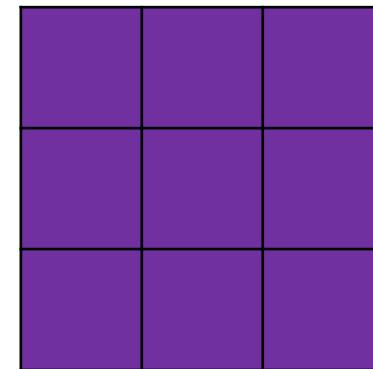
=



\*



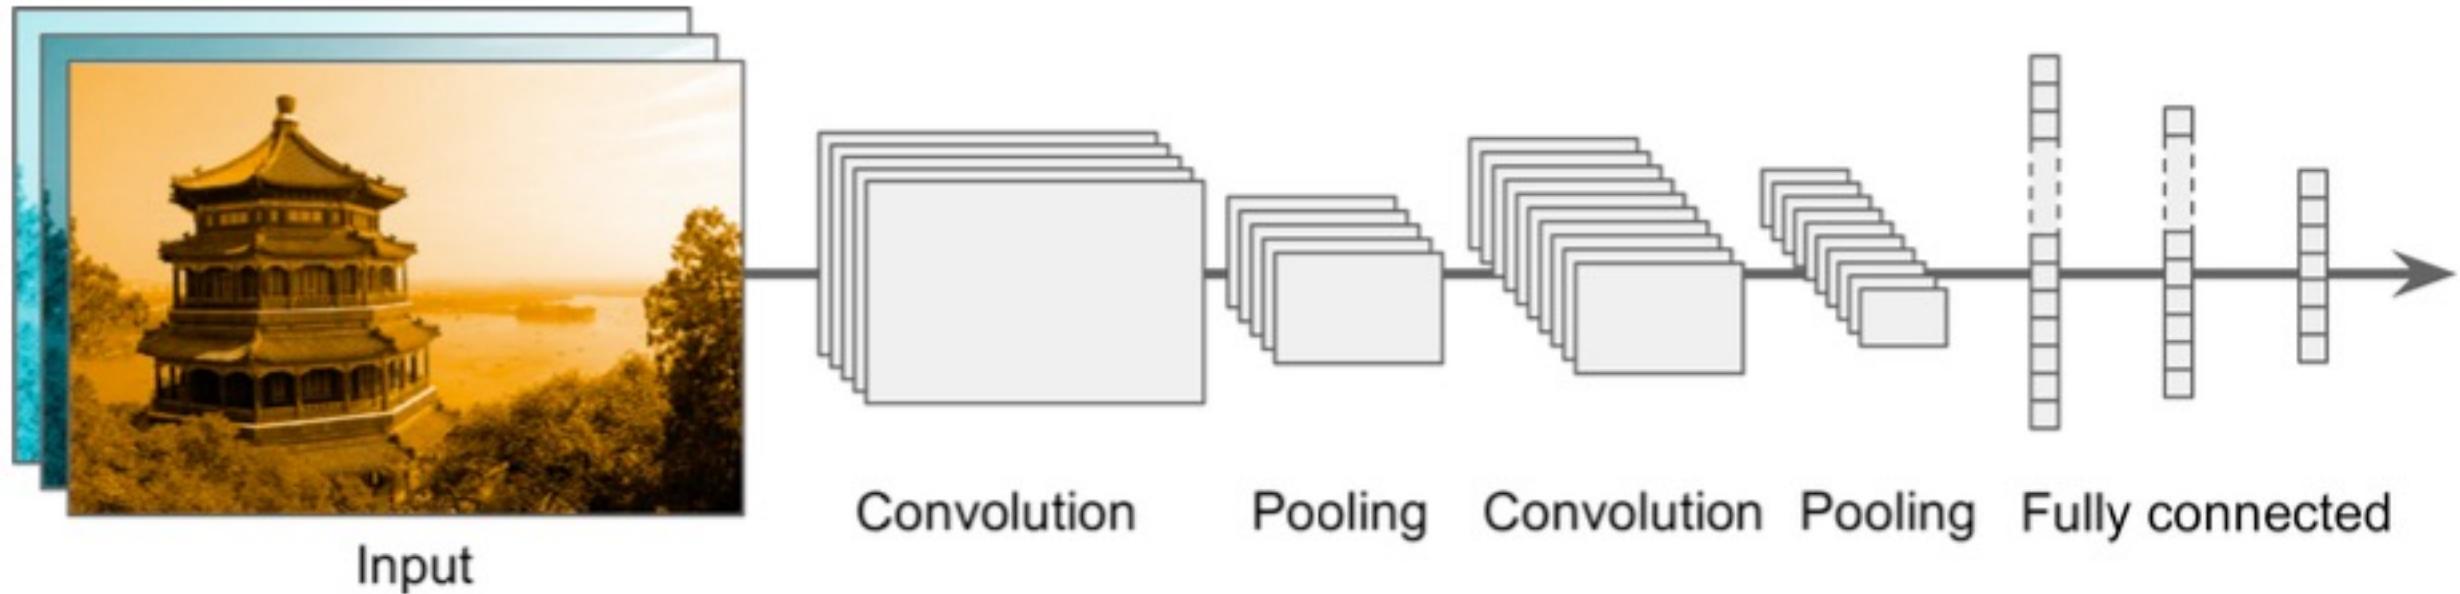
=



$3 \times 3 \times 2$   
 $(n_H^{[1]}, n_W^{[1]}, n_C^{[1]})$

$n_C^{[1]} = \text{number of filters}$

## Typical architecture

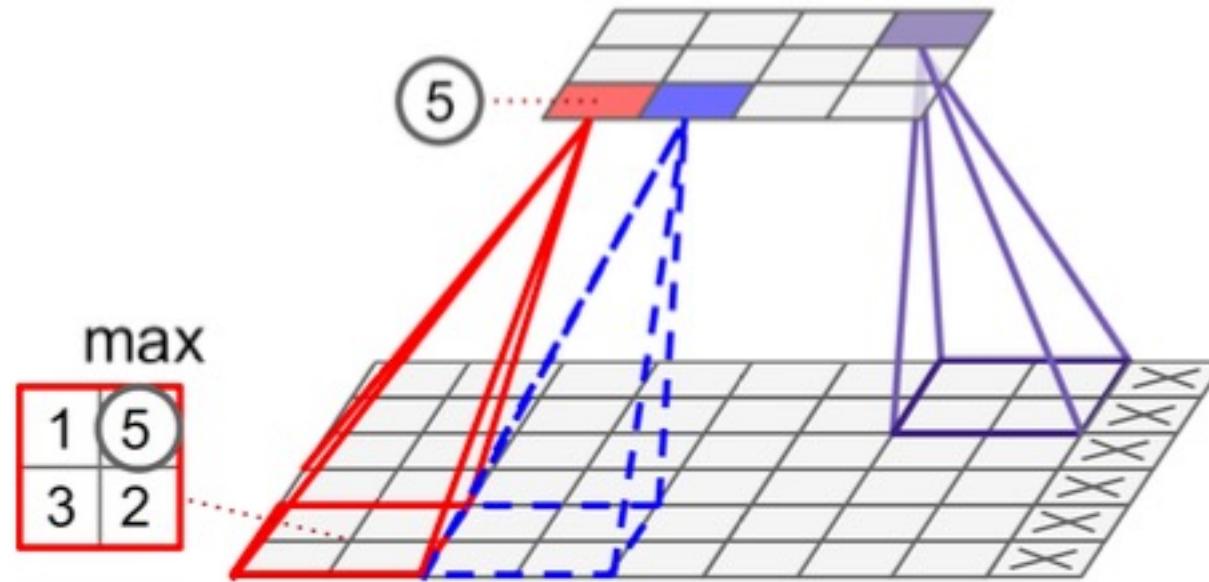


Source: Géron

## Pooling layers

- Subsample (i.e., summarize) the input
  - Reduced computational load
  - Reduced memory usage
  - Fewer parameters (and, thus, less overfitting)

## Max pooling

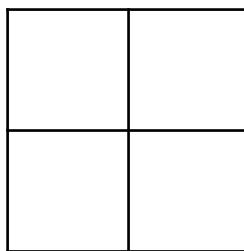


Source: Géron

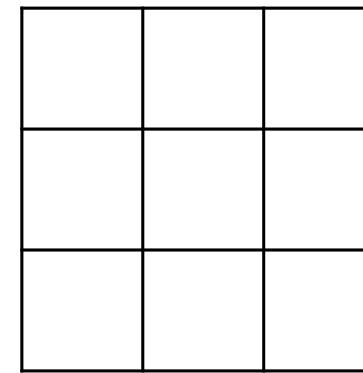
## Max pooling in practice

1	3	1	2
6	1	5	4
5	4	2	5
3	3	1	2

\*



=



Source: Géron

## Max pooling in practice

1	3	1	2
6	1	5	4
5	4	2	5
3	3	1	2

\*


=

6		

Source: Géron

## Max pooling in practice

1	3	1	2
6	1	5	4
5	4	2	5
3	3	1	2

\*


=

6		

Source: Géron

## Max pooling in practice

1	3	1	2
6	1	5	4
5	4	2	5
3	3	1	2

\*


=

6	5	

Source: Géron

## Max pooling in practice

1	3	1	2
6	1	5	4
5	4	2	5
3	3	1	2

\*


=

6	5	

Source: Géron

## Max pooling in practice

1	3	1	2
6	1	5	4
5	4	2	5
3	3	1	2

\*


=

6	5	5

Source: Géron

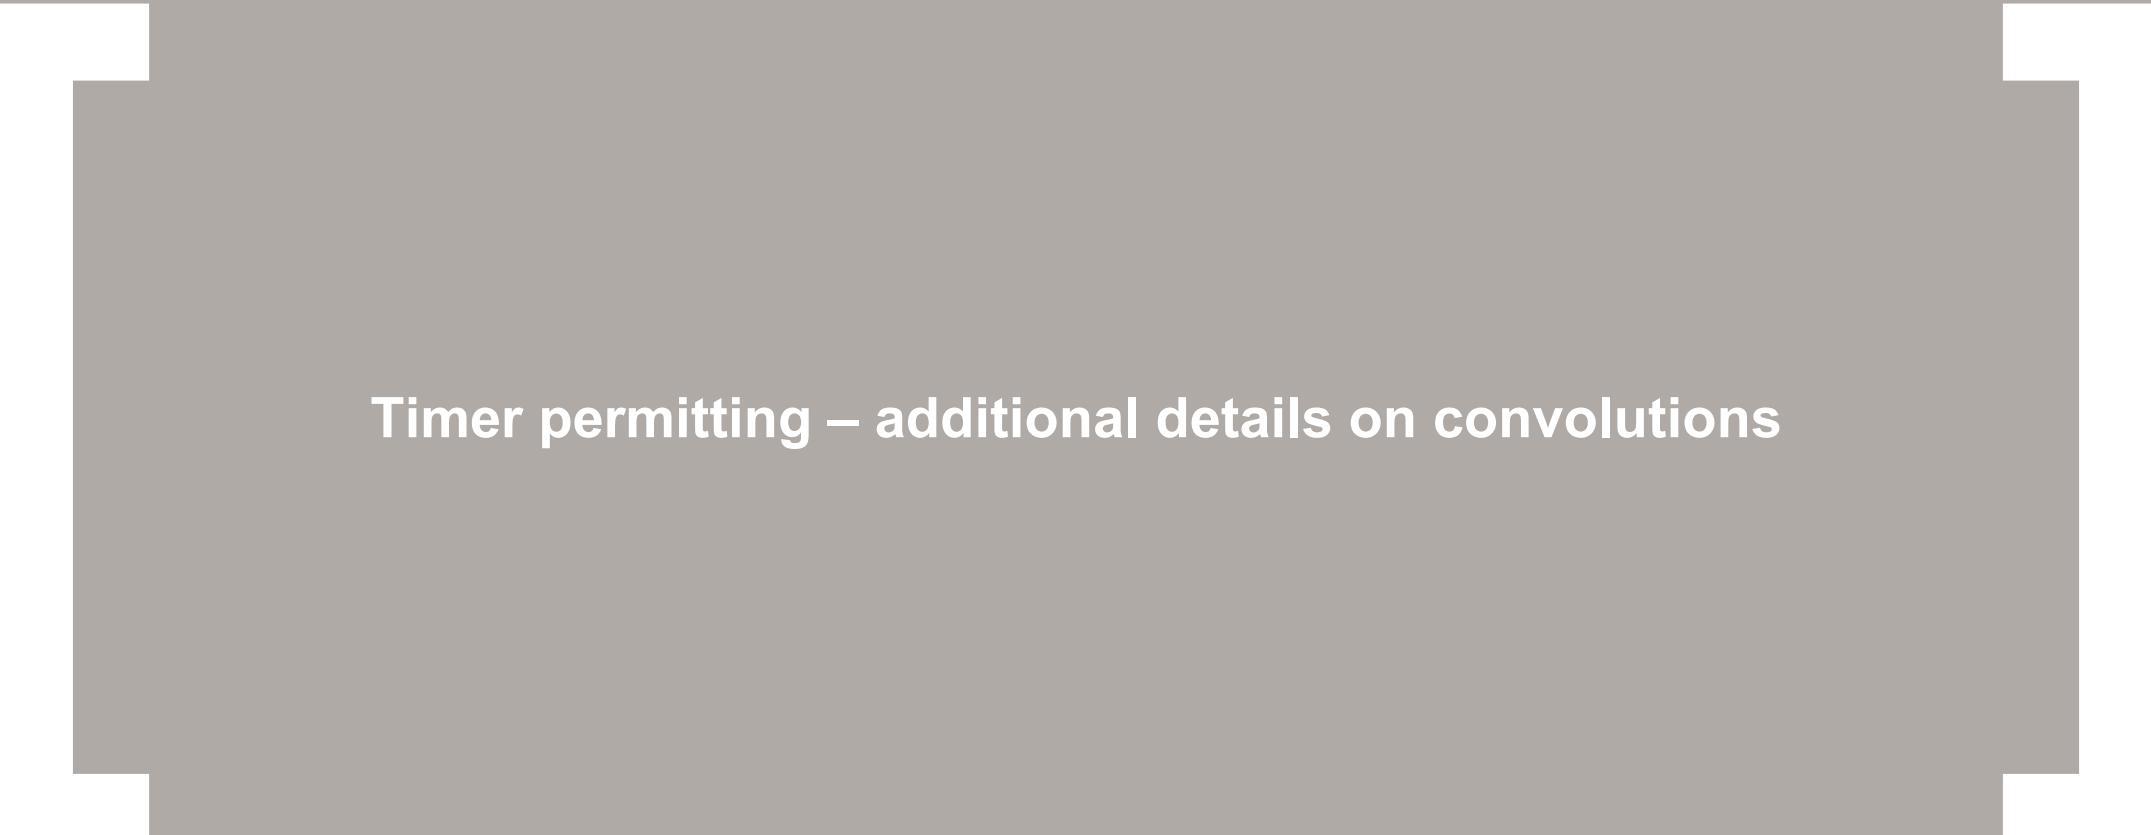
## Max pooling in practice

1	3	1	2
6	1	5	4
5	4	2	5
3	3	1	2

\*


=

6	5	5
6	5	5
5	4	5



**Timer permitting – additional details on convolutions**

## With convolution, the center matters more

1	3	1	2
6	1	5	4
5	4	2	5
3	3	1	2

\*

1	2
2	1

=

20	12	19
22	21	22
22	15	16

## The convolution operator with padding

	1	3	1	2	
	6	1	5	4	
	5	4	2	5	
	3	3	1	2	

\*

1	2
2	1

=


## The convolution operator with padding

	1	3	1	2	
	6	1	5	4	
	5	4	2	5	
	3	3	1	2	

\*

1	2
2	1

=

1				

## The convolution operator with padding

	1	3	1	2	
	6	1	5	4	
	5	4	2	5	
	3	3	1	2	

\*

1	2
2	1

=

1				

## The convolution operator with padding

	1	3	1	2	
	6	1	5	4	
	5	4	2	5	
	3	3	1	2	

\*

1	2
2	1

=

1	4			

## The convolution operator with padding

	1	3	1	2	
	6	1	5	4	
	5	4	2	5	
	3	3	1	2	

\*

1	2
2	1

=

1	4			

## The convolution operator with padding

	1	3	1	2	
	6	1	5	4	
	5	4	2	5	
	3	3	1	2	

\*

1	2
2	1

=

1	4	7		

## The convolution operator with padding

	1	3	1	2	
	6	1	5	4	
	5	4	2	5	
	3	3	1	2	

\*

1	2
2	1

=

1	4	7	4	4
8	20	12	19	10
17	22	21	22	14
13	22	15	16	9
6	9	5	5	2

## The convolution operator with padding

	1	3	1	2	
	6	1	5	4	
	5	4	2	5	
	3	3	1	2	

\*

1	2
2	1

=

1	4	7	4	4
8	20	12	19	10
17	22	21	22	14
13	22	15	16	9
6	9	5	5	2

## The convolution operator with padding and stride

	1	3	1	2	
	6	1	5	4	
	5	4	2	5	
	3	3	1	2	

\*

1	2
2	1

=


## The convolution operator with padding and stride

	1	3	1	2	
	6	1	5	4	
	5	4	2	5	
	3	3	1	2	

\*

1	2
2	1

=

1		

## The convolution operator with padding and stride

	1	3	1	2	
	6	1	5	4	
	5	4	2	5	
	3	3	1	2	

\*

1	2
2	1

=

1		

## The convolution operator with padding and stride

	1	3	1	2	
	6	1	5	4	
	5	4	2	5	
	3	3	1	2	

\*

1	2
2	1

=

1	7	

## The convolution operator with padding and stride

	1	3	1	2	
	6	1	5	4	
	5	4	2	5	
	3	3	1	2	

\*

1	2
2	1

=

1	7	

## The convolution operator with padding and stride

	1	3	1	2	
	6	1	5	4	
	5	4	2	5	
	3	3	1	2	

\*

1	2
2	1

=

1	7	4

## The convolution operator with padding and stride

	1	3	1	2		
	6	1	5	4		
	5	4	2	5		
	3	3	1	2		

\*

1	2
2	1

=

1	7	4
17	21	14
6	5	2

## The convolution operator with padding and stride

	1	3	1	2	
	6	1	5	4	
	5	4	2	5	
	3	3	1	2	

4x4  
 $(n_H^{[0]}, n_W^{[0]})$

\*  
2x2 with padding 1x1 and stride 2x2  
 $(f_H, f_W)$  with  $(p_H, p_W)$  and  $(s_H, s_W)$

1	2
2	1

=

1	7	4
17	21	14
6	5	2

3x3  
 $(n_H^{[1]}, n_W^{[1]})$

## The convolution operator with padding and stride

	1	3	1	2		
	6	1	5	4		
	5	4	2	5		
	3	3	1	2		

4x4  
 $(n_H^{[0]}, n_W^{[0]})$

2x2 with padding 1x1 and stride 2x2  
 $(f_H, f_W)$  with  $(p_H, p_W)$  and  $(s_H, s_W)$

$$n_H^{[1]} = \left\lfloor \frac{n_H^{[0]} + 2p_H - f_H}{s_H} + 1 \right\rfloor \quad n_W^{[1]} = \left\lfloor \frac{n_W^{[0]} + 2p_W - f_W}{s_W} + 1 \right\rfloor$$

1	2
2	1

=

1	7	4
17	21	14
6	5	2



BAYES  
 BUSINESS SCHOOL  
CITY UNIVERSITY OF LONDON



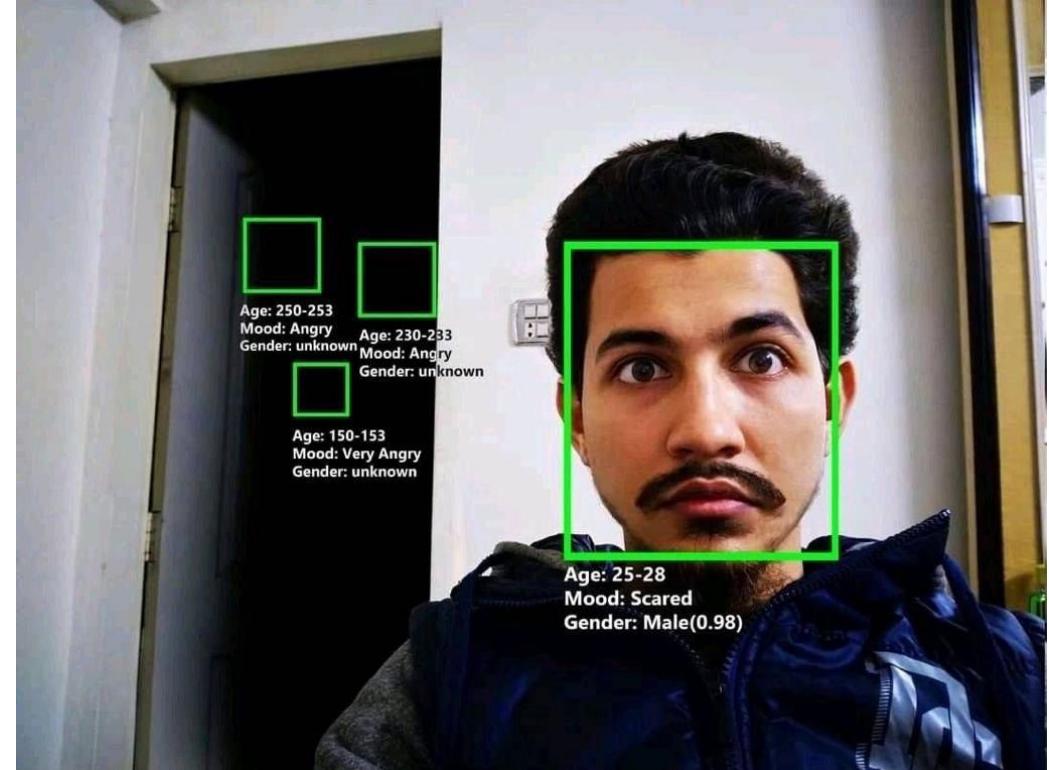
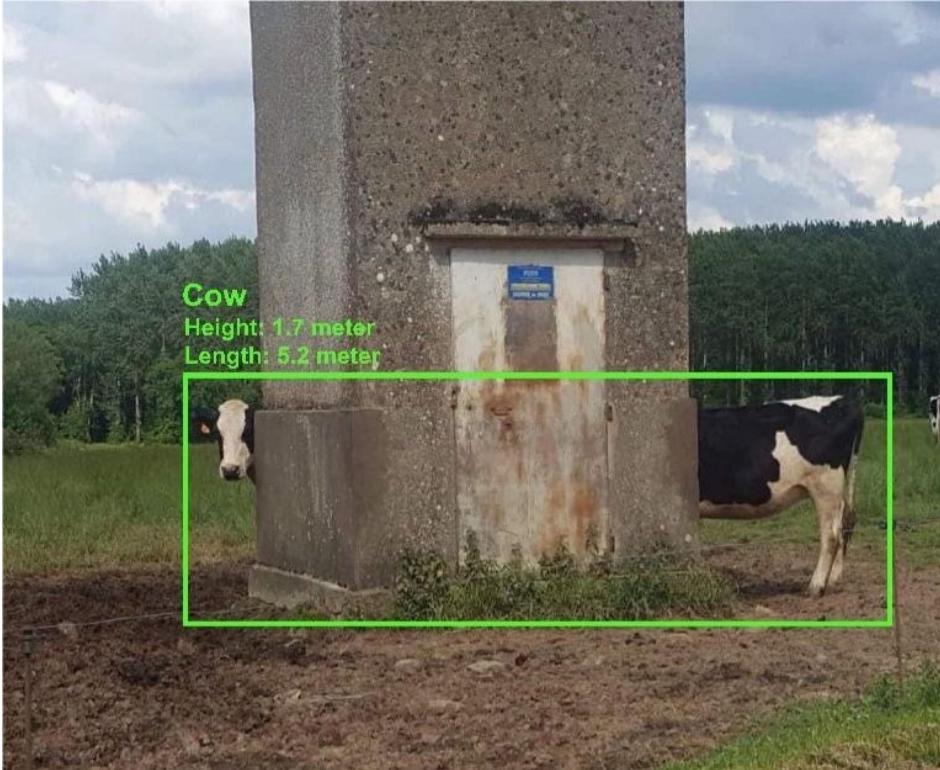
**Image classification as a business**

# Use cases and challenges



**Beyond image classification**

# Object detection



## Semantic segmentation



Source: Towards AI, 2021

# Image generation



DALL·E

History

Collections

Edit the detailed description

Surprise me

Upload →

A lecture theatre full of students in business analytics

Generate



**BAYES**  
BUSINESS SCHOOL  
CITY UNIVERSITY OF LONDON

Source: OpenAI



See you next week!



## Sources

- Alla & Kalyan Adari, 2019, Beginning Anomaly Detection Using Python-Based Deep Learning, Apress
- Bhaskhar, 2021, Introduction to Deep Learning: <https://cs229.stanford.edu/syllabus.html>
- Chollet, 2021, Deep Learning with Python, 2<sup>nd</sup> edition
- DeepLearning.AI, n.d.: [deeplearning.ai](https://deeplearning.ai)
- Dieleman, 2020, Lecture 3: Convolutional Neural Networks:  
[https://storage.googleapis.com/deepmind-media/UCLxDeepMind\\_2020/L3%20-%20UUCLxDeepMind%20DL2020.pdf](https://storage.googleapis.com/deepmind-media/UCLxDeepMind_2020/L3%20-%20UUCLxDeepMind%20DL2020.pdf)
- Géron, 2019, Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow
- Goodfellow, Bengio, Courville, 2016, The Deep Learning Book:  
<http://www.deeplearningbook.org>
- Liang, 2016, Introduction to Deep Learning:  
<https://www.cs.princeton.edu/courses/archive/spring16/cos495/>
- Lin et al., 2014, Microsoft COCO: Common Objects in Context:  
<https://arxiv.org/pdf/1405.0312.pdf>
- OpenAI et al., n.d., DALL-E 2: <https://openai.com/dall-e-2/>
- Towards AI, 2021, Semantic Segmentation: A Complete Guide:  
<https://towardsai.net/p/l/machine-learning-7>